

# ArtEdge

## Real-Time Neural Style Transfer (NST) for Mobile

Nayeem  
Mohammad

Nitin  
Gupta

# Project Overview

Neural style transfer (NST) is a computer vision technique that applies the **artistic style of one image to the content of another**.

*ArtEdge* develops a mobile-first solution for **real-time neural style transfer**, allowing users to transform ordinary photos into their artistic variants directly on their mobile devices.

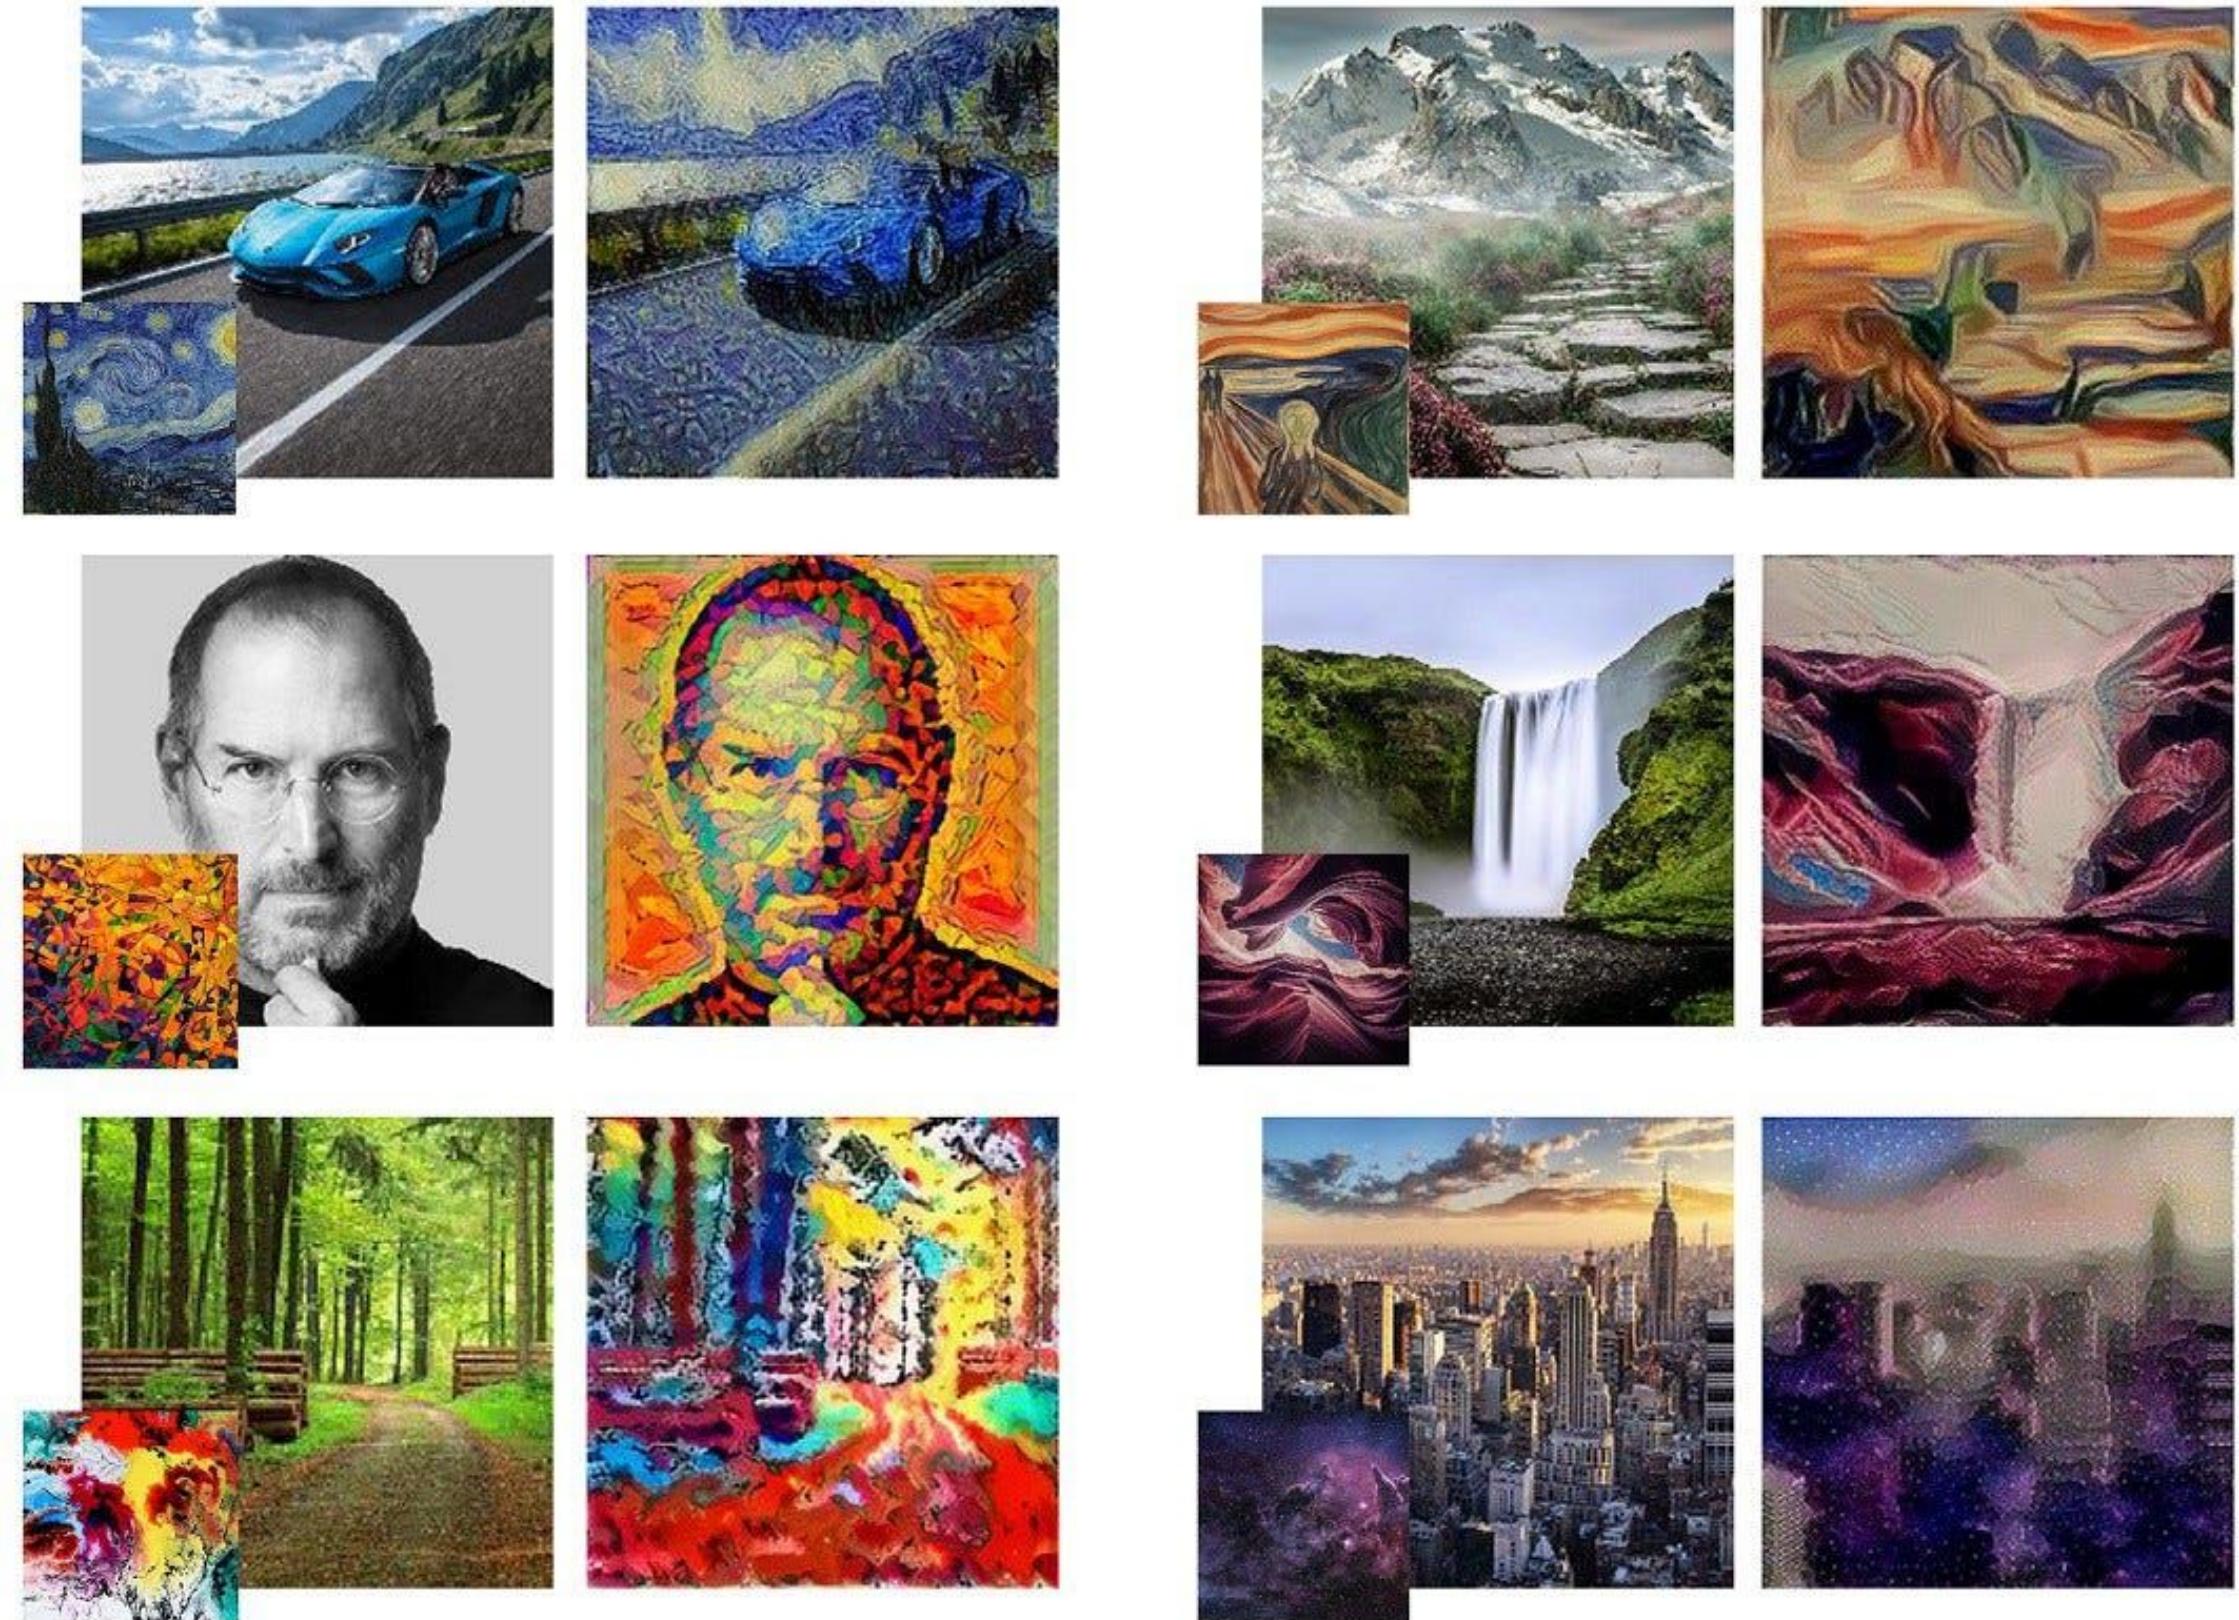


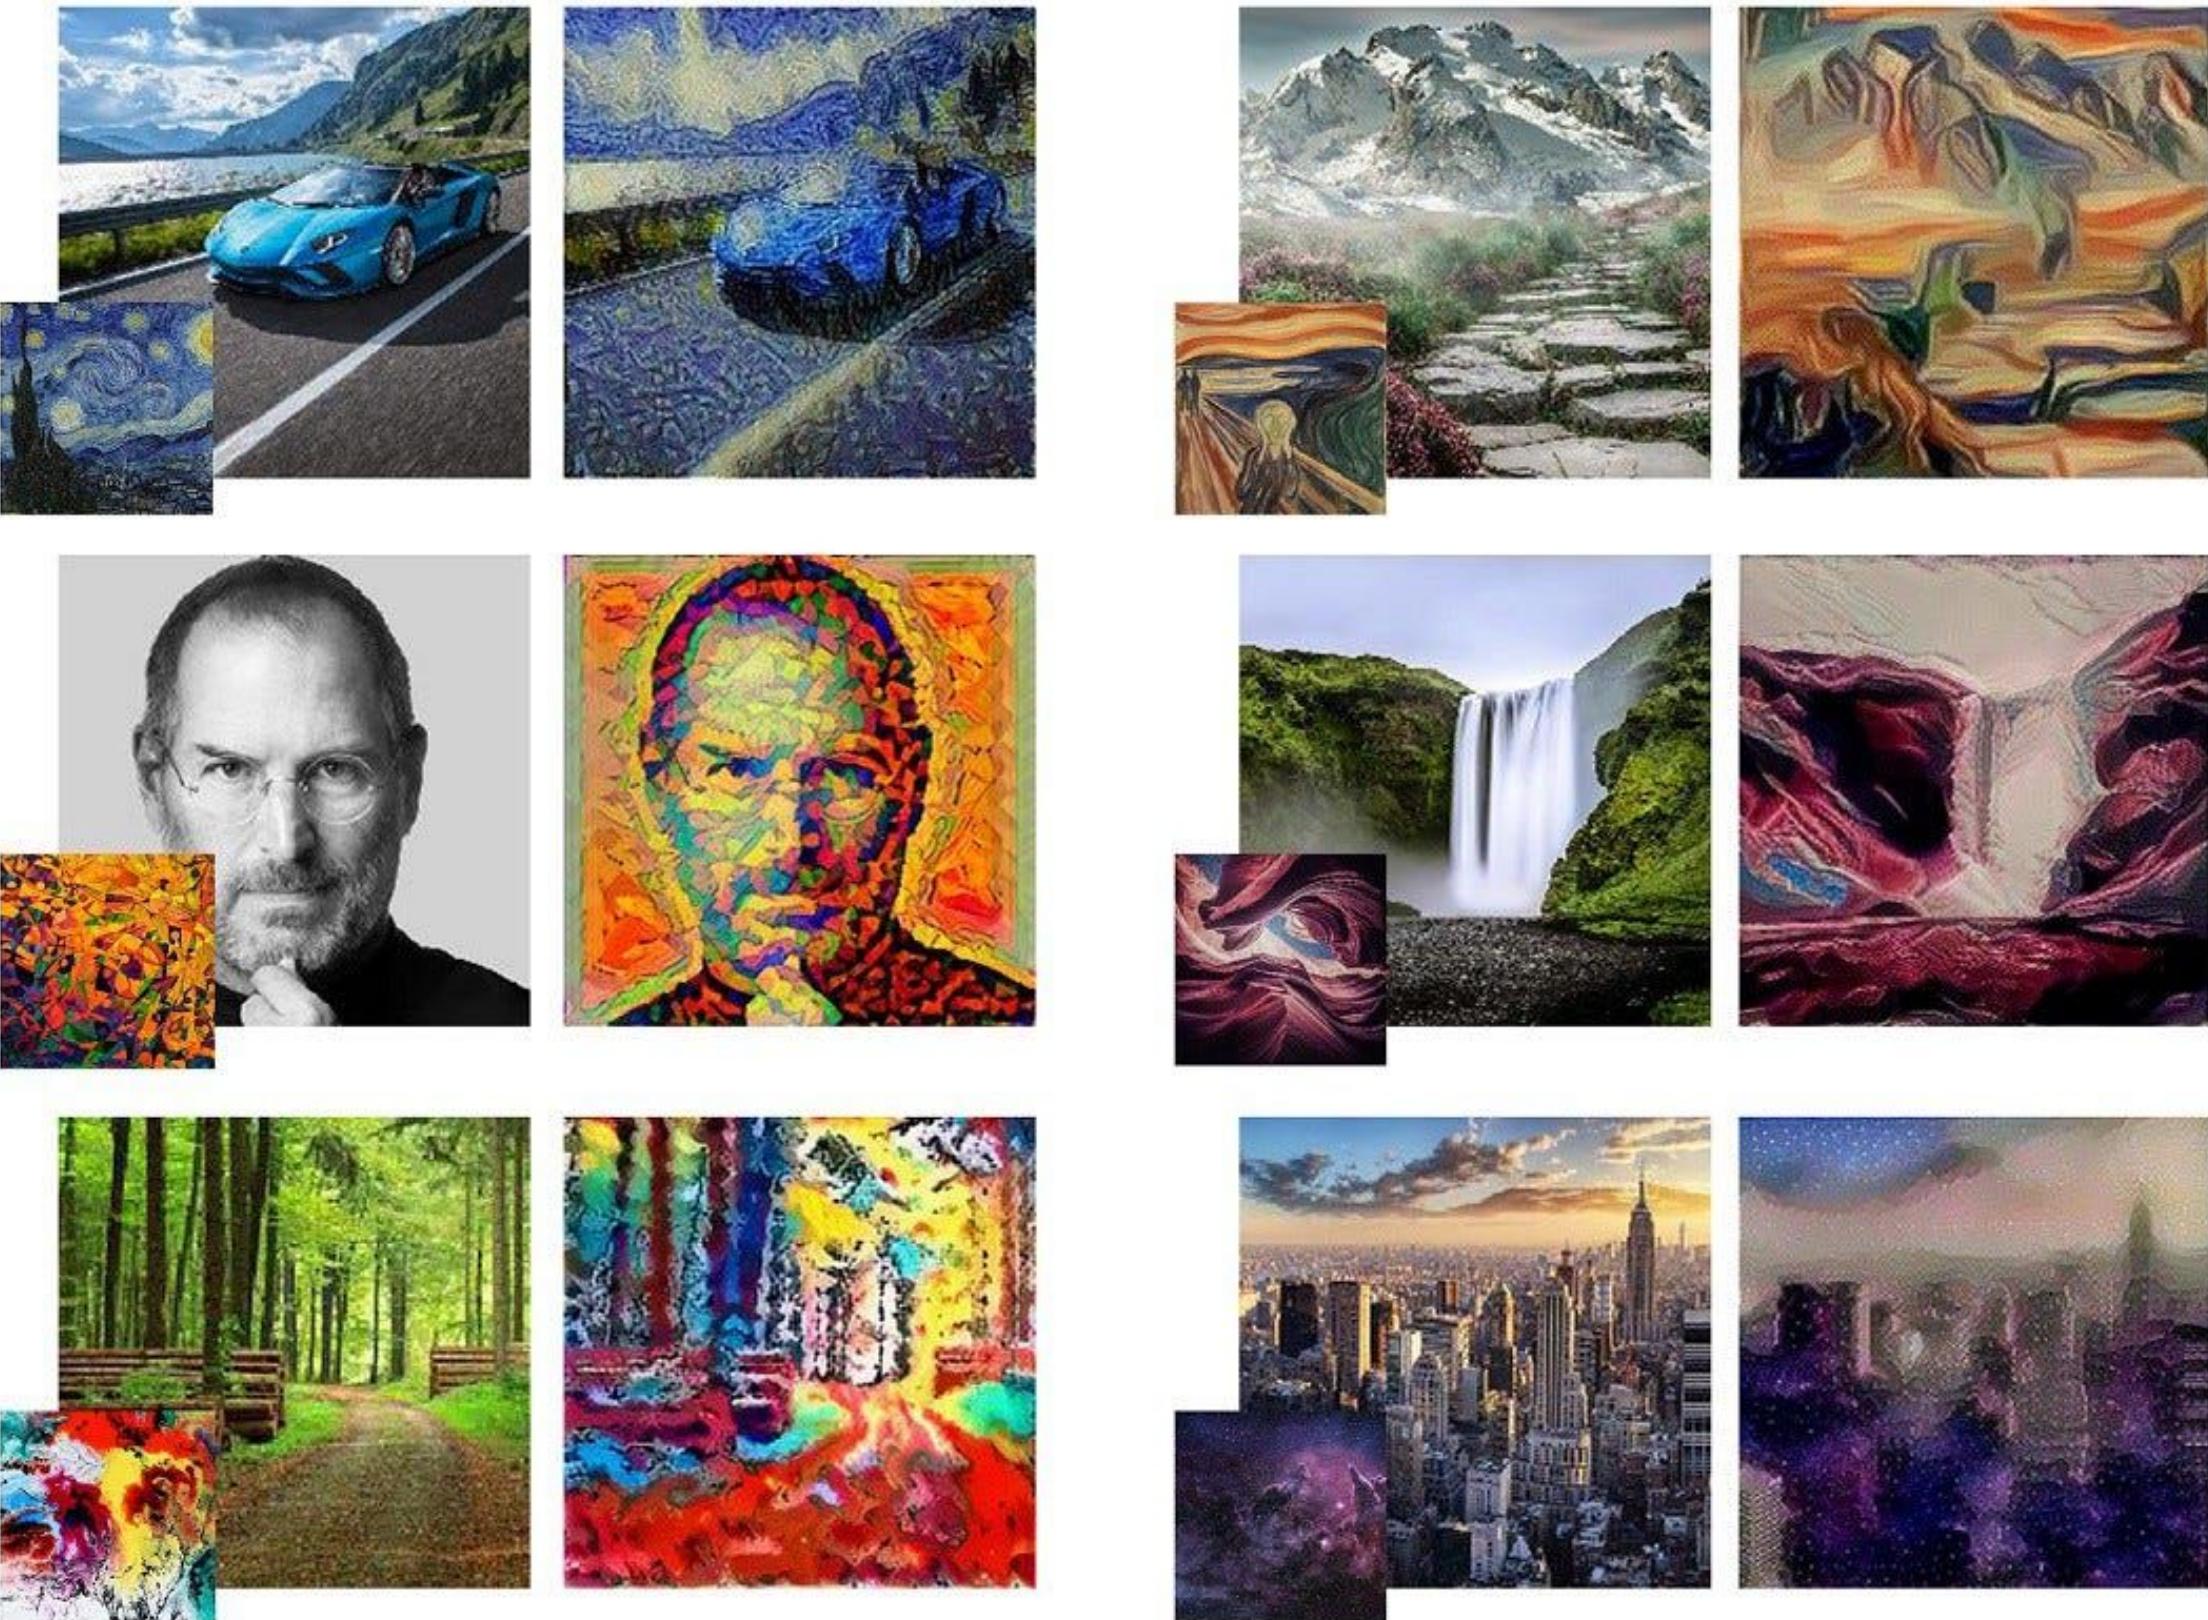
Image Source: [Creatives need neural "concept" style transfer](#)

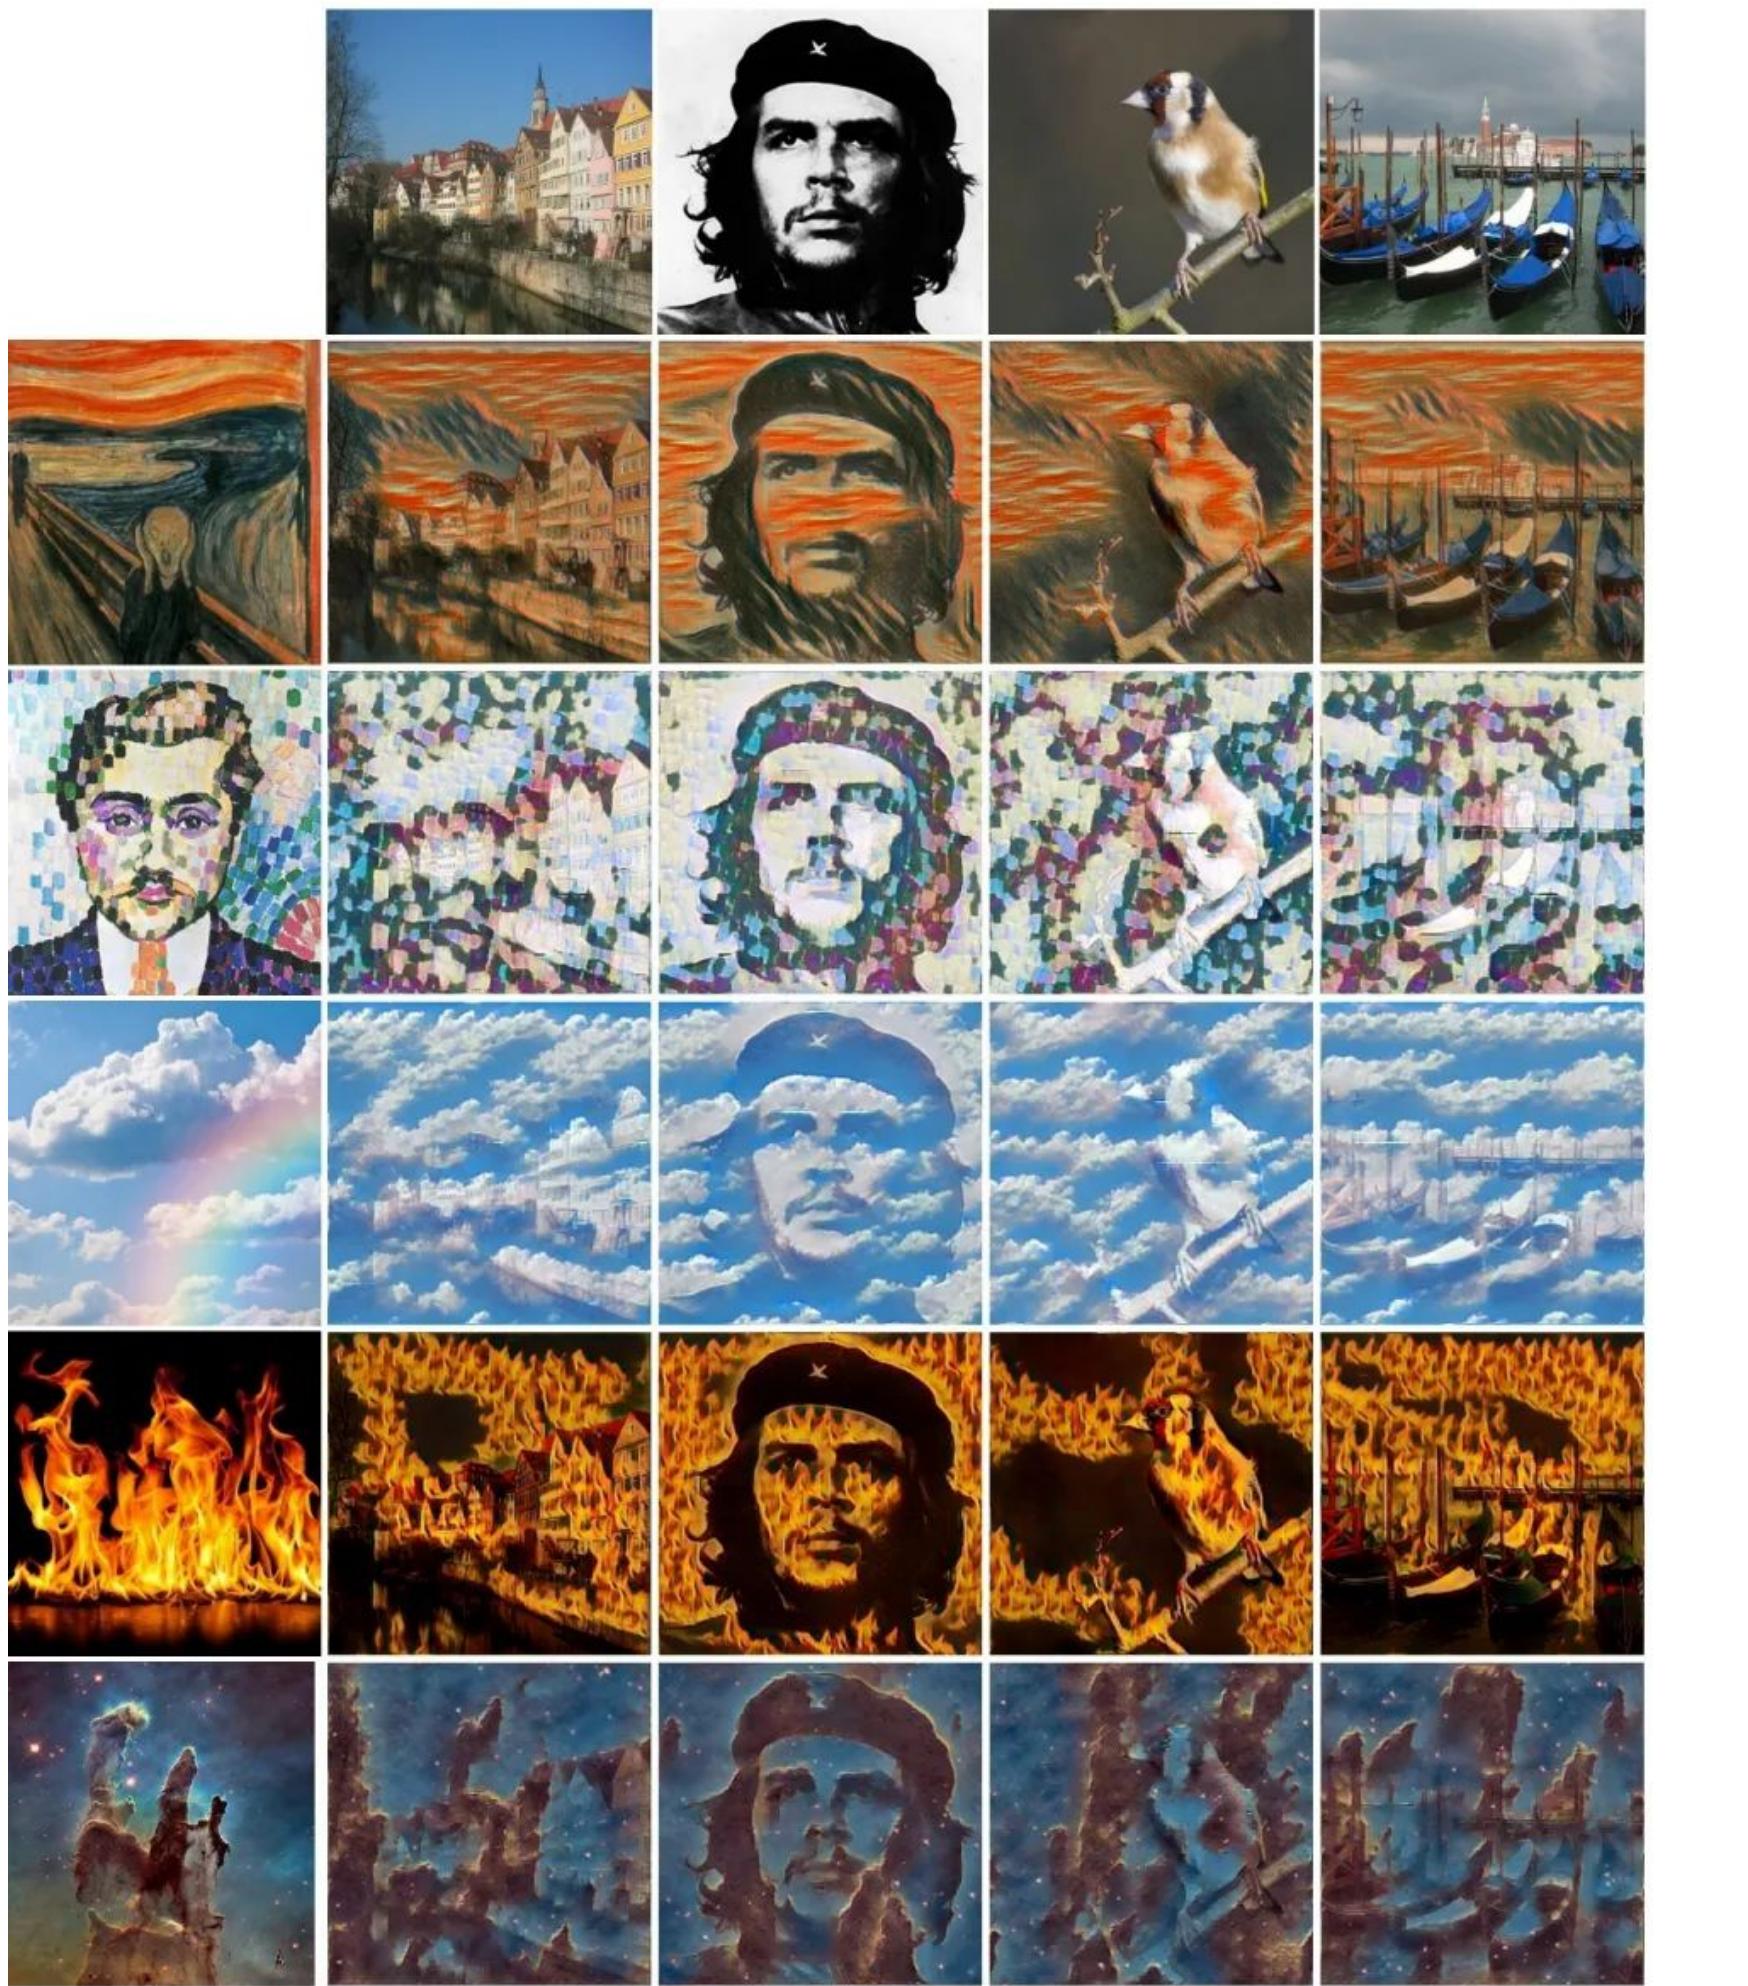
# Project Importance

Enhanced privacy by **processing sensitive images locally** rather than in the cloud

**Democratization of artistic tools**, making sophisticated AI art accessible to everyday users

Creating **more inclusive visual communication** technologies





# Edge Computing

## Relevance

Moving computation from the cloud to the device,  
**reducing latency** from seconds to milliseconds

# Enhancing privacy by keeping personal photos on-device

enabling **offline functionality** for users without reliable connectivity

**developing domain-agnostic techniques,  
applicable to other edge AI applications**

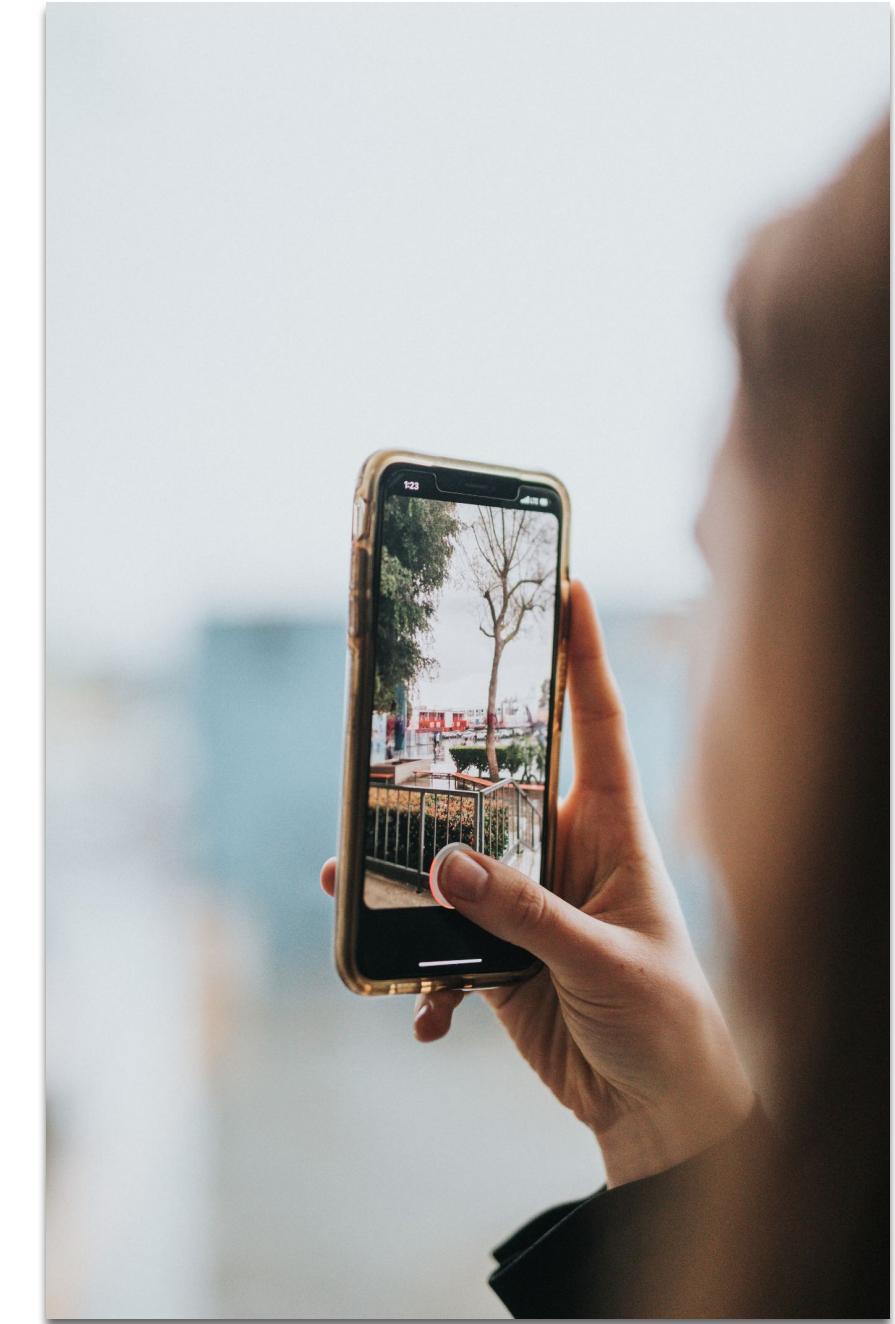
# Datasets



**WikiArt: ~80,000 fine art paintings as 2D color images** with variable resolutions (ranging from **500 to 3000 px**), featuring diverse artistic styles.



**MS COCO: ~330,000 2D color images containing 1.5 million object instances** with variable resolutions (typically around **640×480 px**), featuring diverse scenes, objects, and environments.



# User-Generated Images

Custom **real-world 2D color** (RGB) photos captured on mobile devices features variable resolutions (typically **1080×1920 px**), representing diverse lighting conditions and subjects.

# Literature Review

# 1. Original Neural Style Transfer

**Challenge:** Maintaining original image content while applying a completely different artistic aesthetic

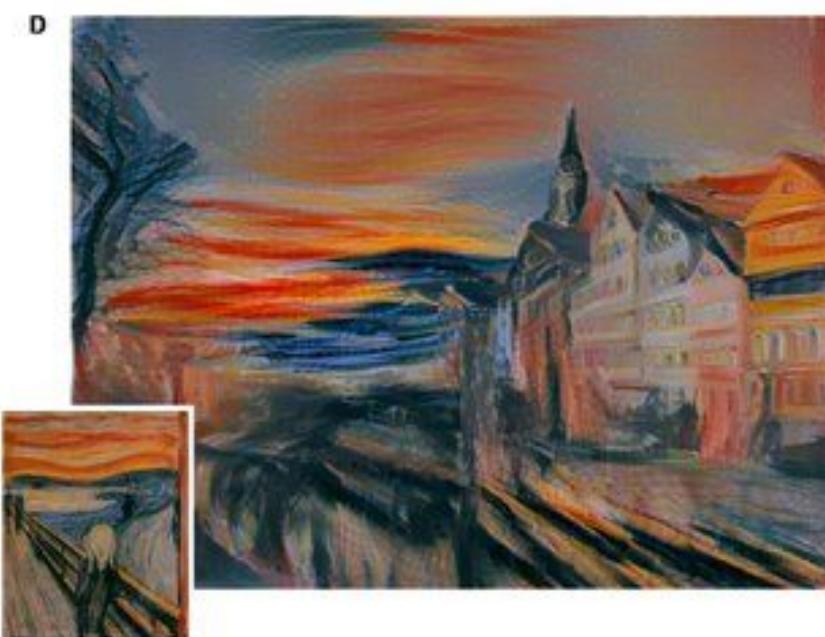
**Insights:**

Different CNN layers capture different visual information

Lower layers: Texture and local patterns

Higher layers: Semantic content and global structure

Style is more than just color - it's about feature correlations and distributions



Source: [A Neural Algorithm of Artistic Style \(Gatys et al., 2016\)](#)

# 1. Original Neural Style Transfer

Used **pre-trained VGG-19** network

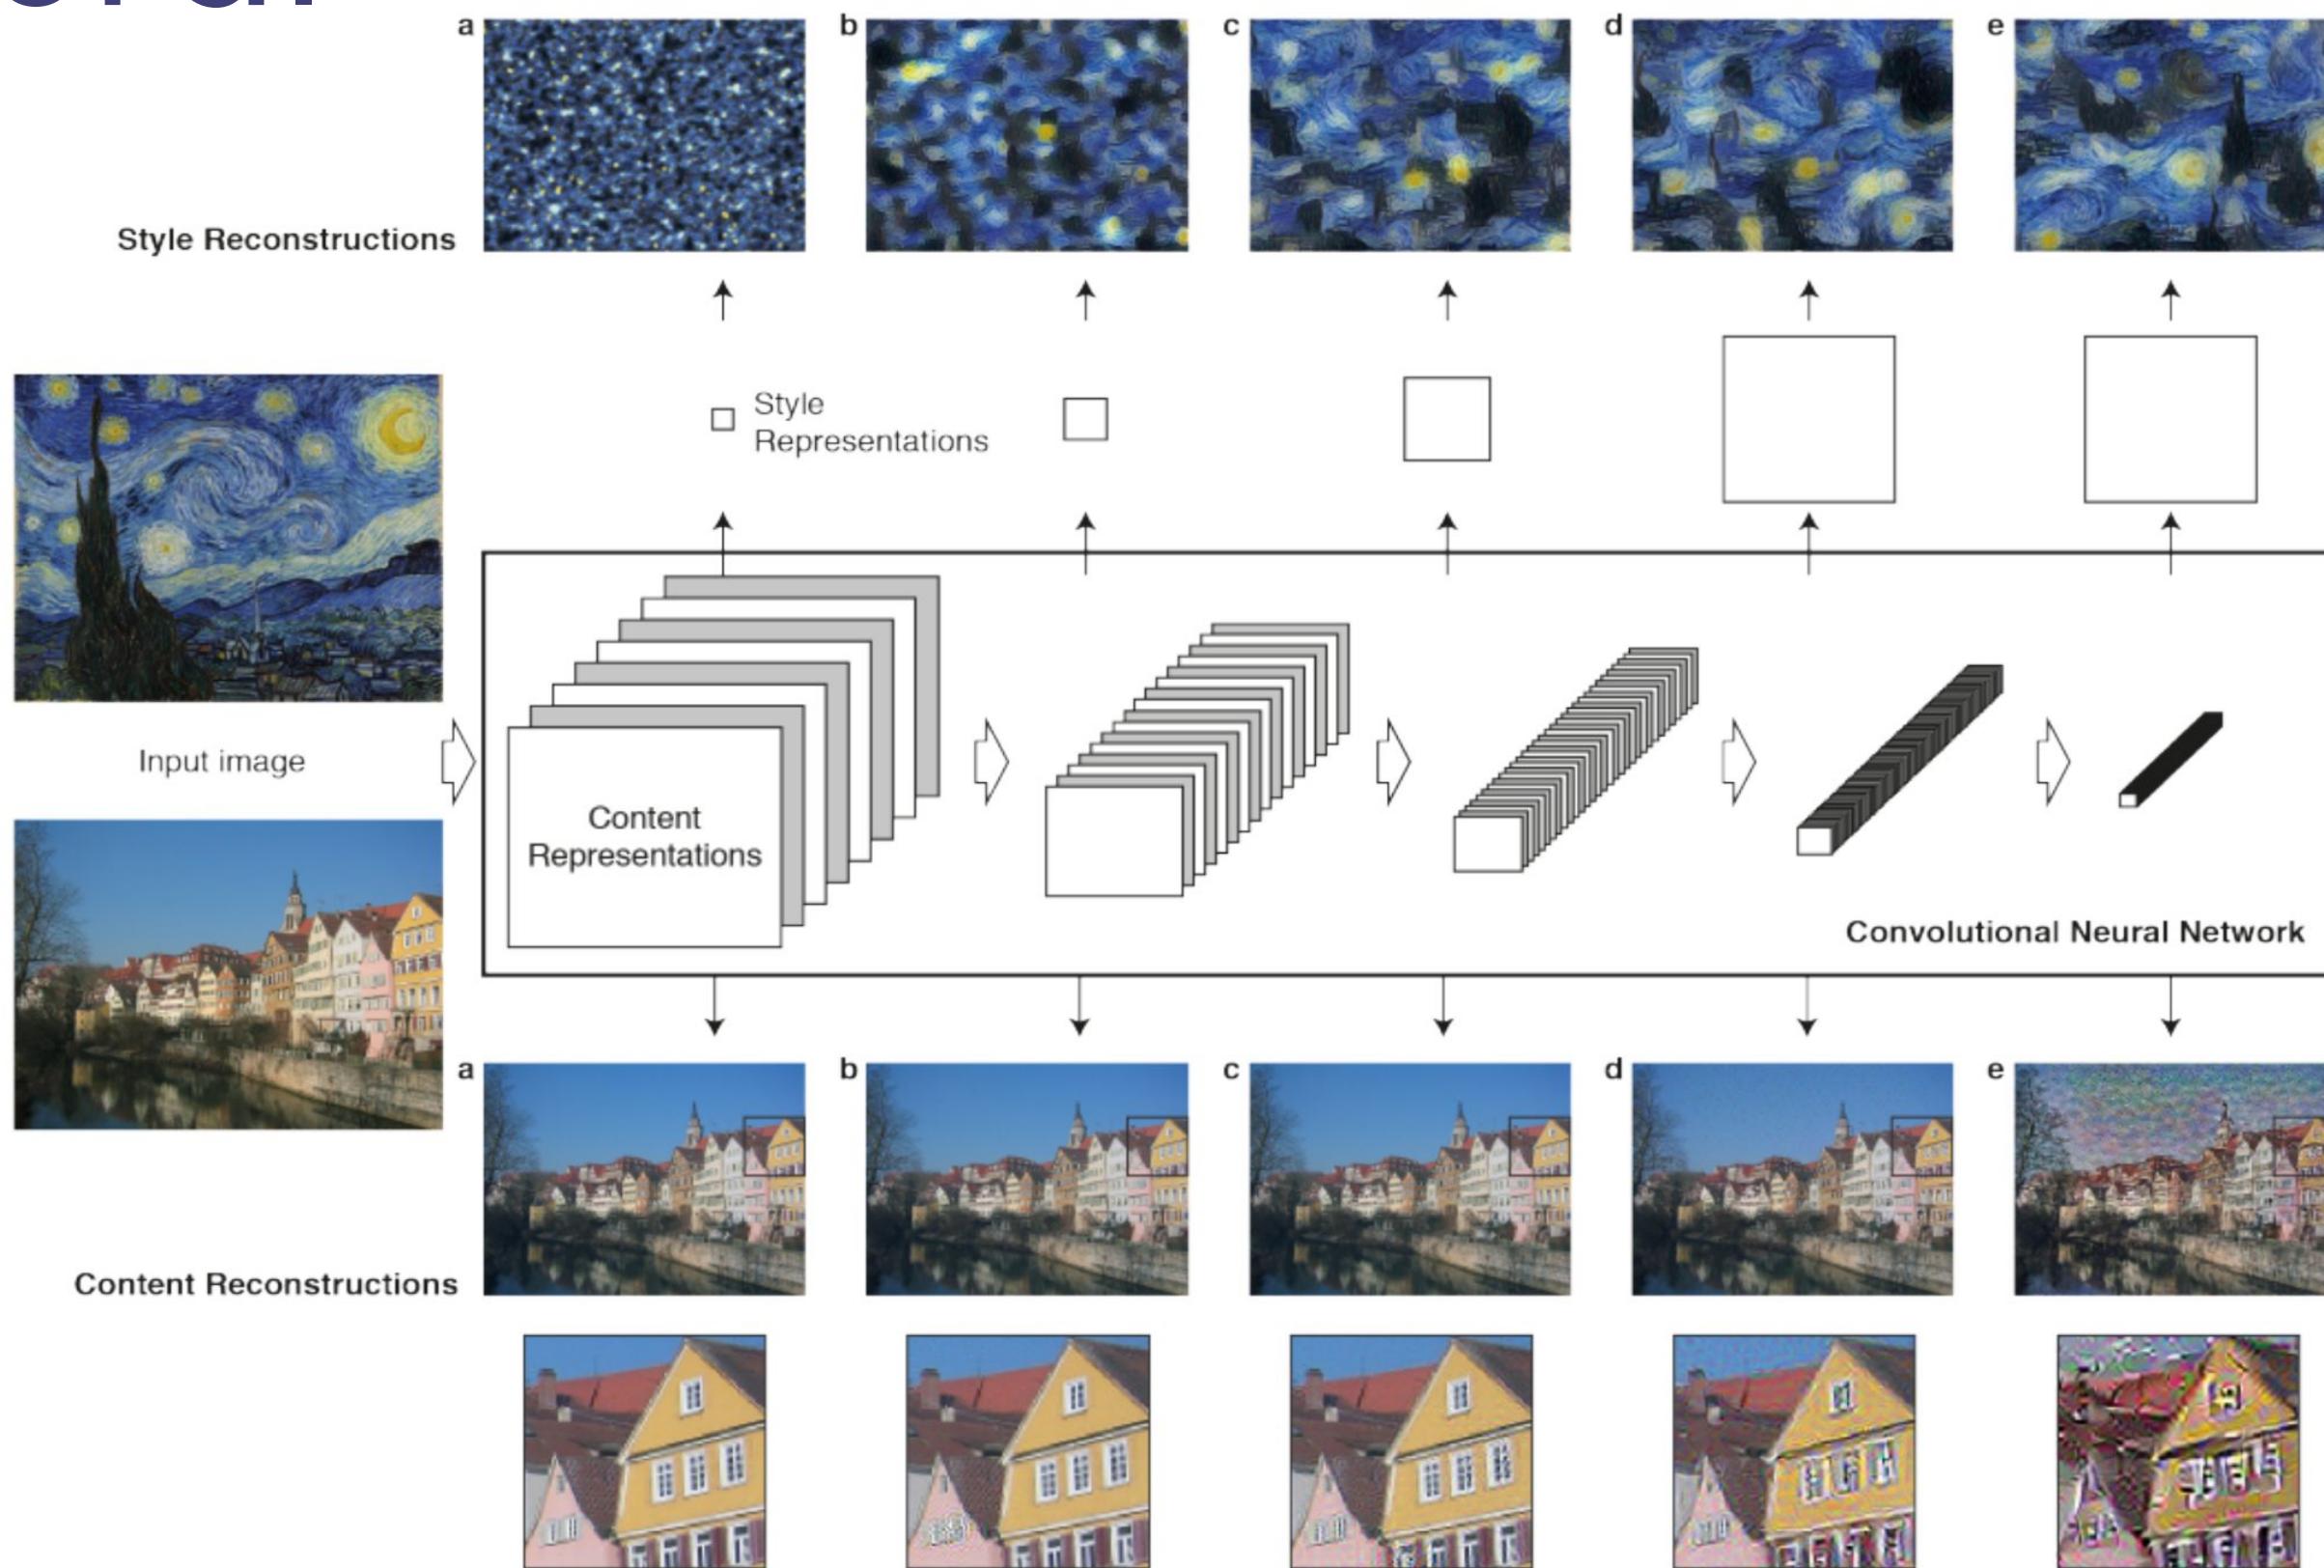
Leveraged features from multiple layers

Loss is a **weighted combination** of content and style losses

**Iteratively updates** input image

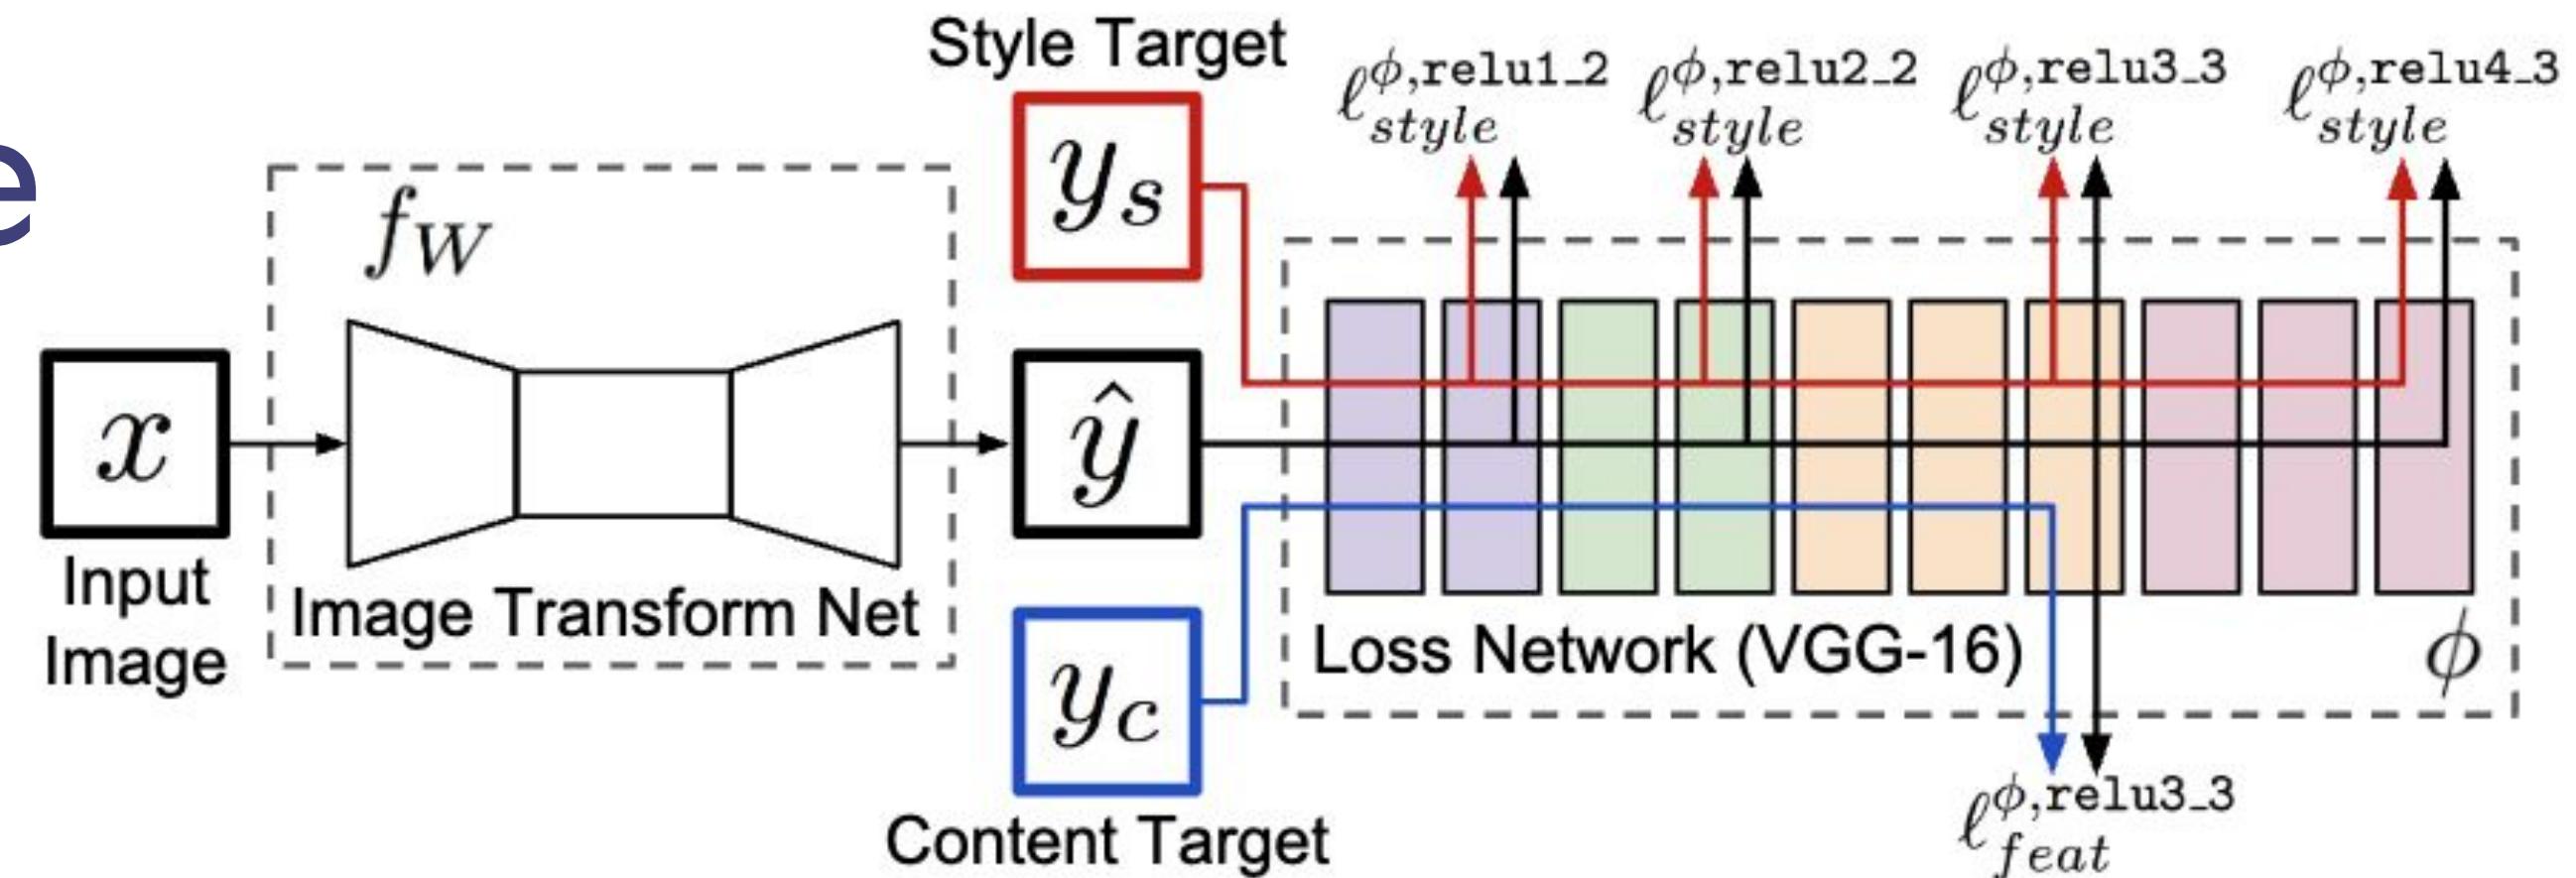
Worked across multiple artistic styles

Requires **significant computational resources**



Source: [A Neural Algorithm of Artistic Style \(Gatys et al., 2016\)](#)

## 2. Fast Style Transfer



Replaces traditional per-pixel loss with **perceptual loss** (based on high-level features from a pretrained CNN) to train feed-forward networks for Neural Style Transfer.

**Training:** Single-Style  
(train one network per style image)

### Perceptual Loss:

**Feature Reconstruction Loss:** Ensures content preservation using intermediate VGG features

**Style Reconstruction Loss:** Matches Gram matrices of features to capture textures and style

# 3. Improved NST: AdaIN

Combination of:

1. Flexibility of optimization-based methods
2. The speed of feed-forward approaches.

## Batch Normalization

$$\text{BN}(x) = \gamma \left( \frac{x - \mu(x)}{\sigma(x)} \right) + \beta$$

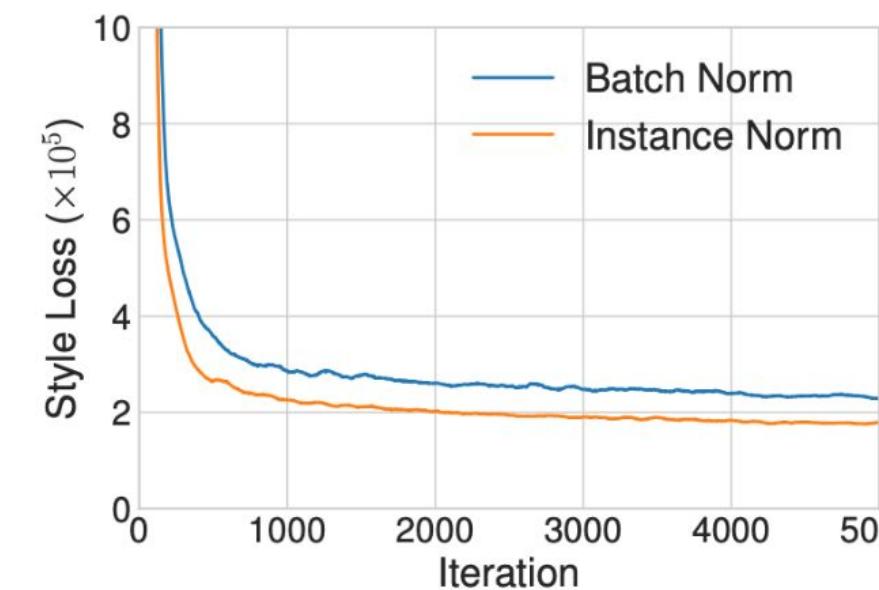
$$\mu_c(x) = \frac{1}{NHW} \sum_{n=1}^N \sum_{h=1}^H \sum_{w=1}^W x_{nchw}$$

## Instance Normalization

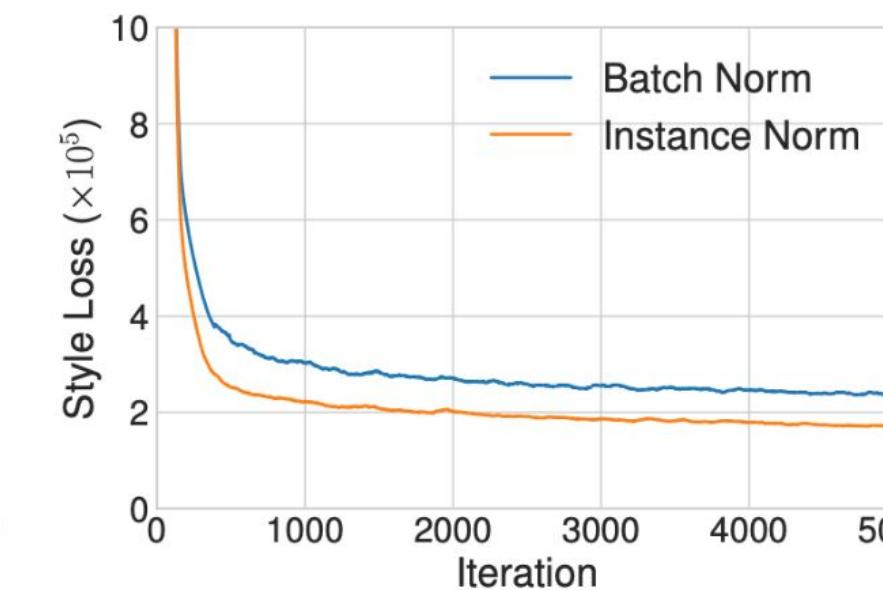
$$\text{IN}(x) = \gamma \left( \frac{x - \mu(x)}{\sigma(x)} \right) + \beta$$

$$\mu_{nc}(x) = \frac{1}{HW} \sum_{h=1}^H \sum_{w=1}^W x_{nchw}$$

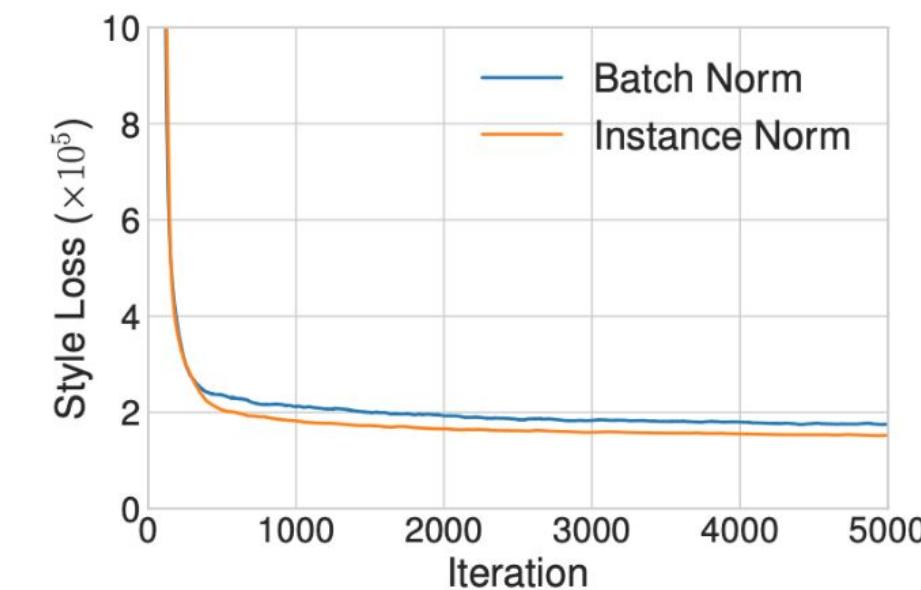
Idea: Adaptive Instance Normalization (AdaIN) simply **adjusts the mean and variance** of the content input to match those of the style input.



(a) Trained with original images.



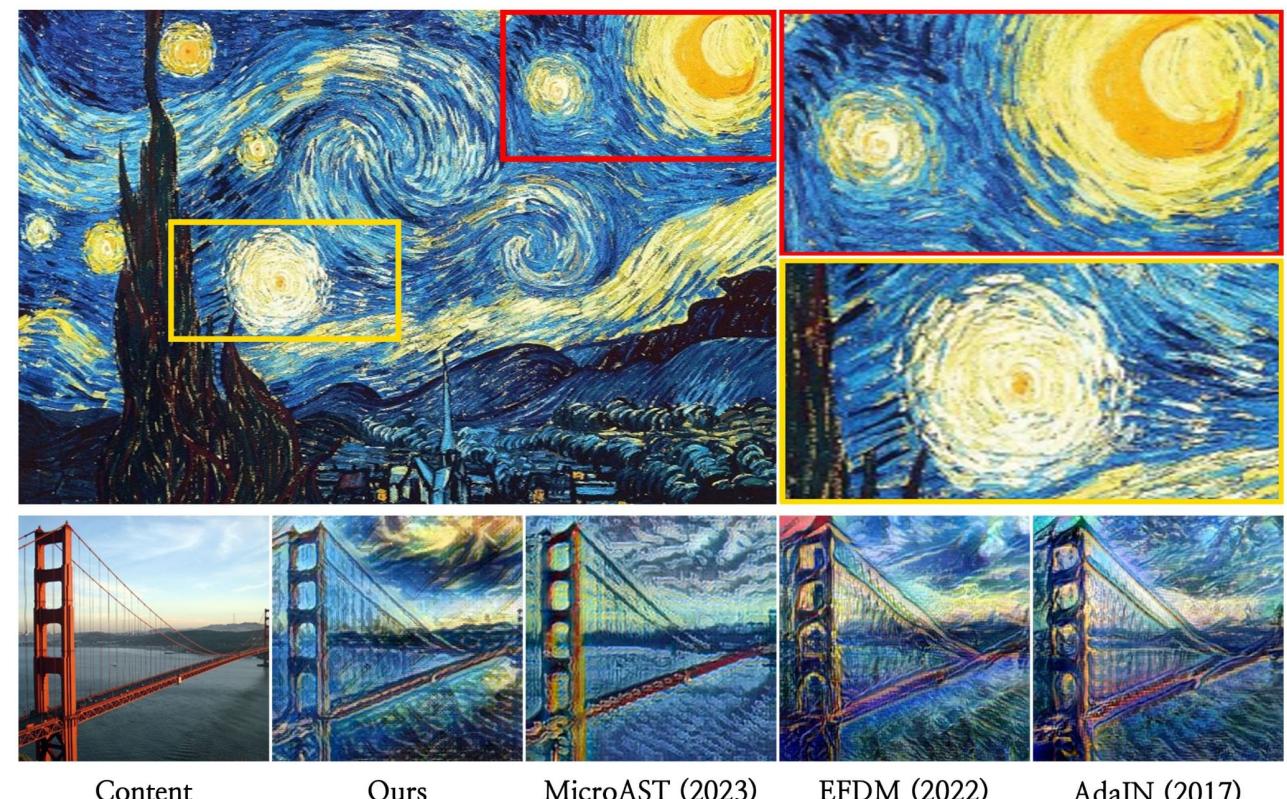
(b) Trained with contrast normalized images.



(c) Trained with style normalized images.

# 4. NST on Edge (AesFA)

Existing methods **lack spatial coherence** and **rely on heavy pre-trained models** making them inefficient for mobile/edge use.



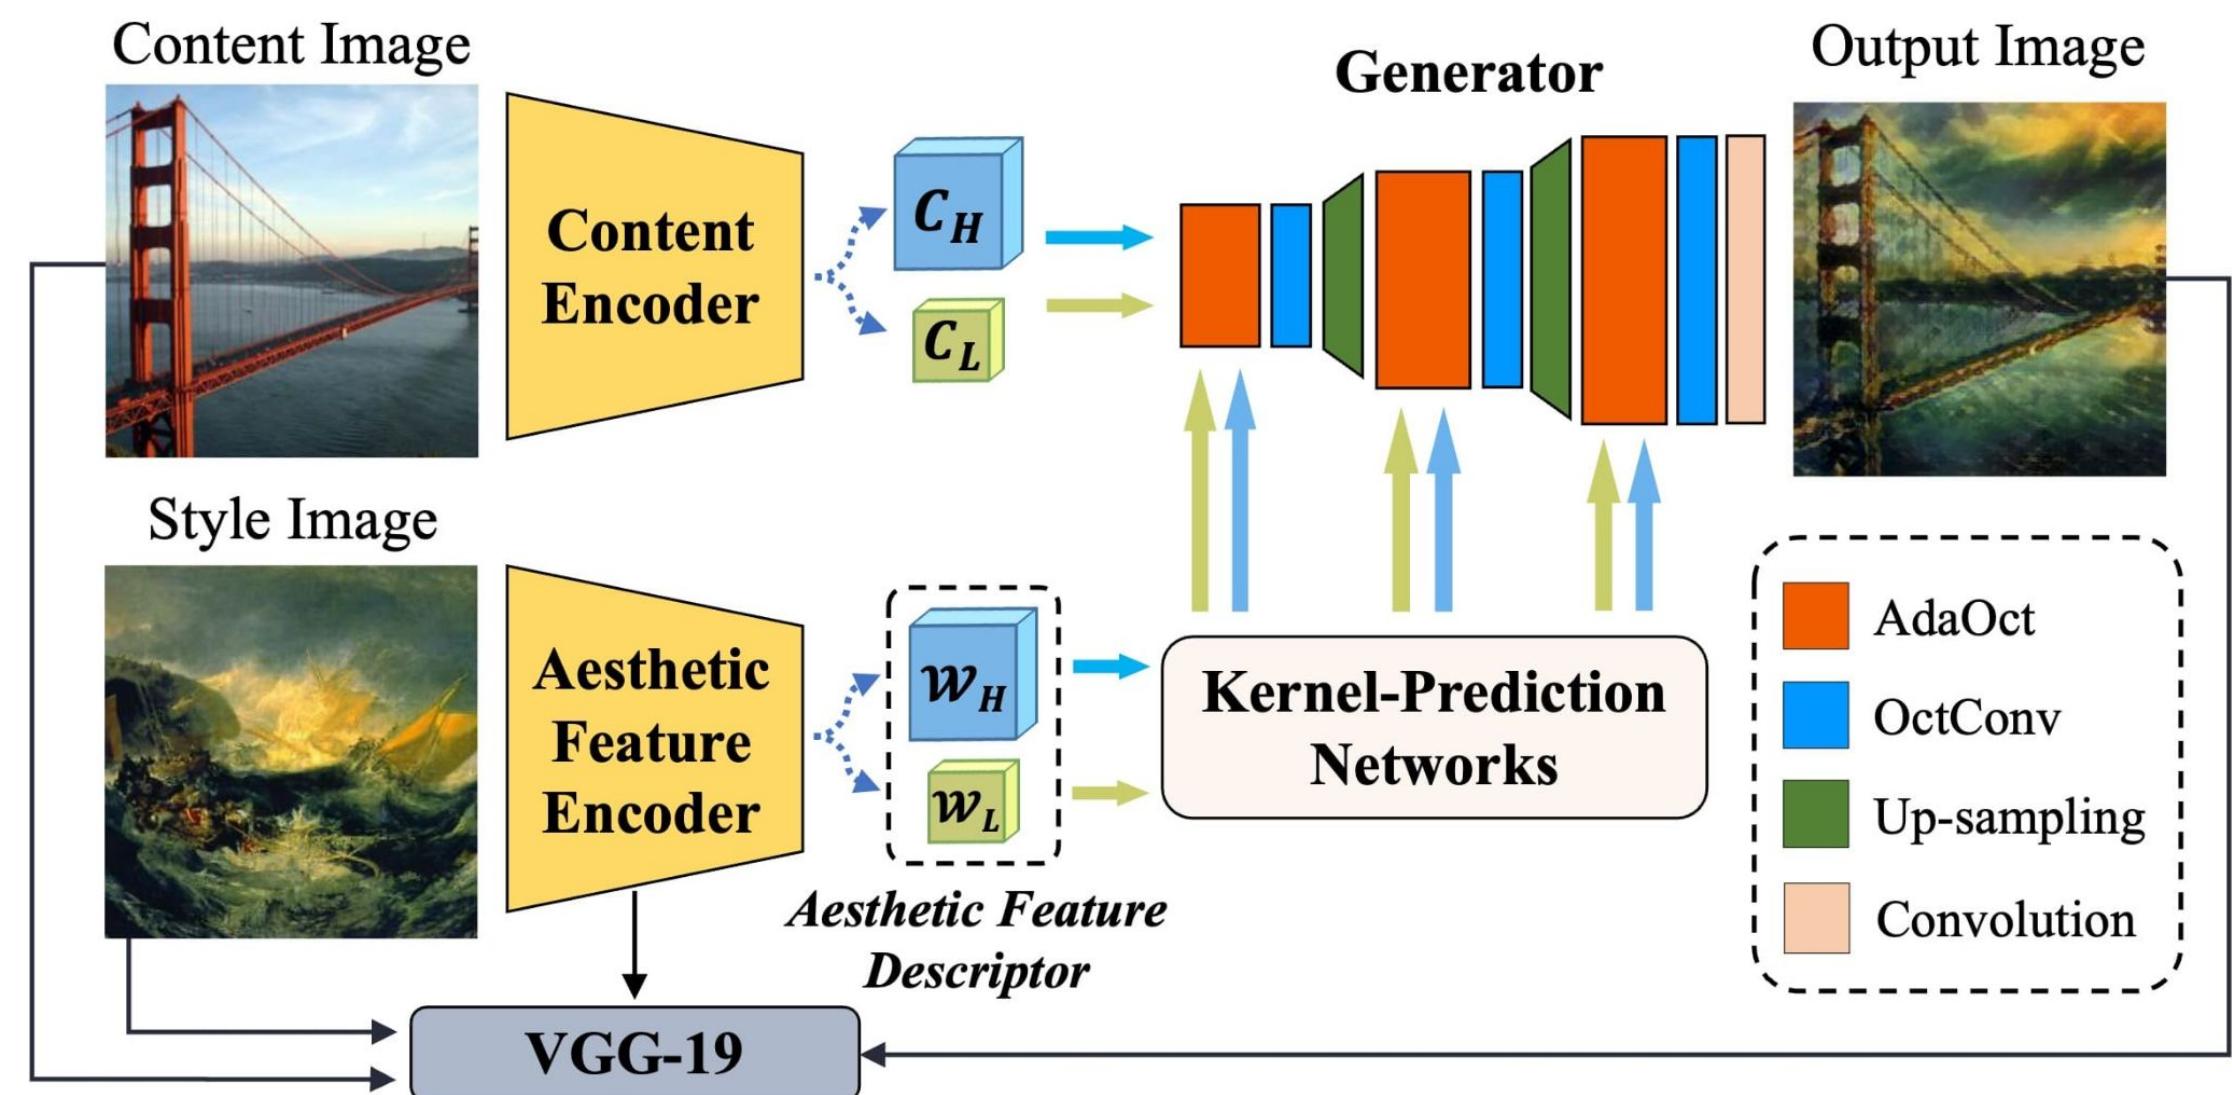
## Innovations:

Splits into **high-frequency** (textures/brushstrokes) and **low-frequency** (global tones)

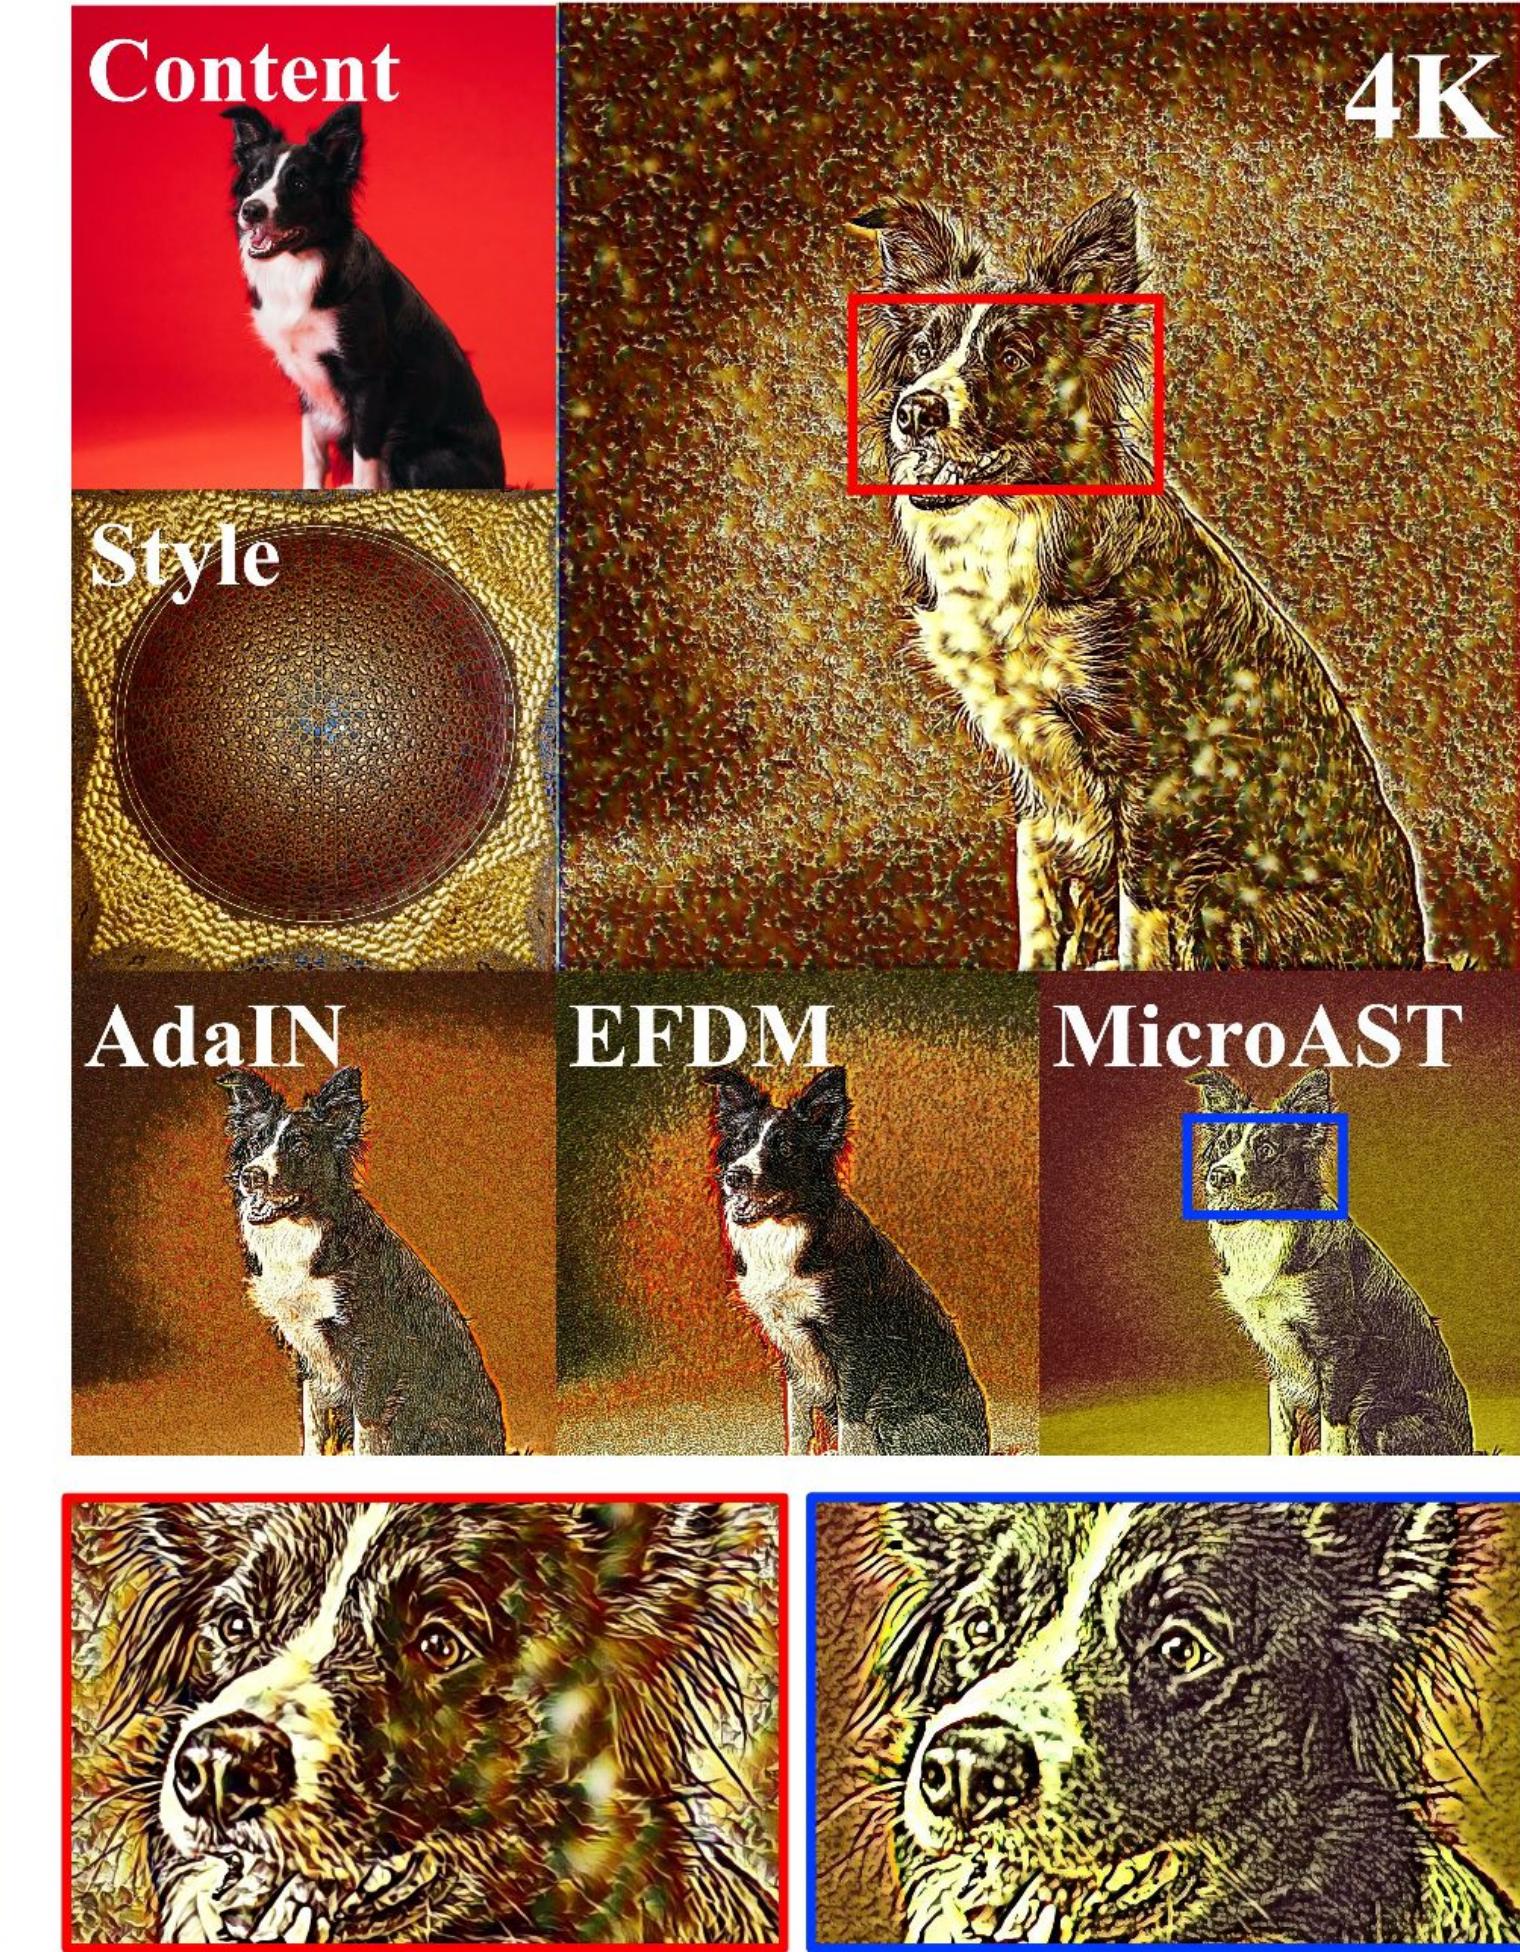
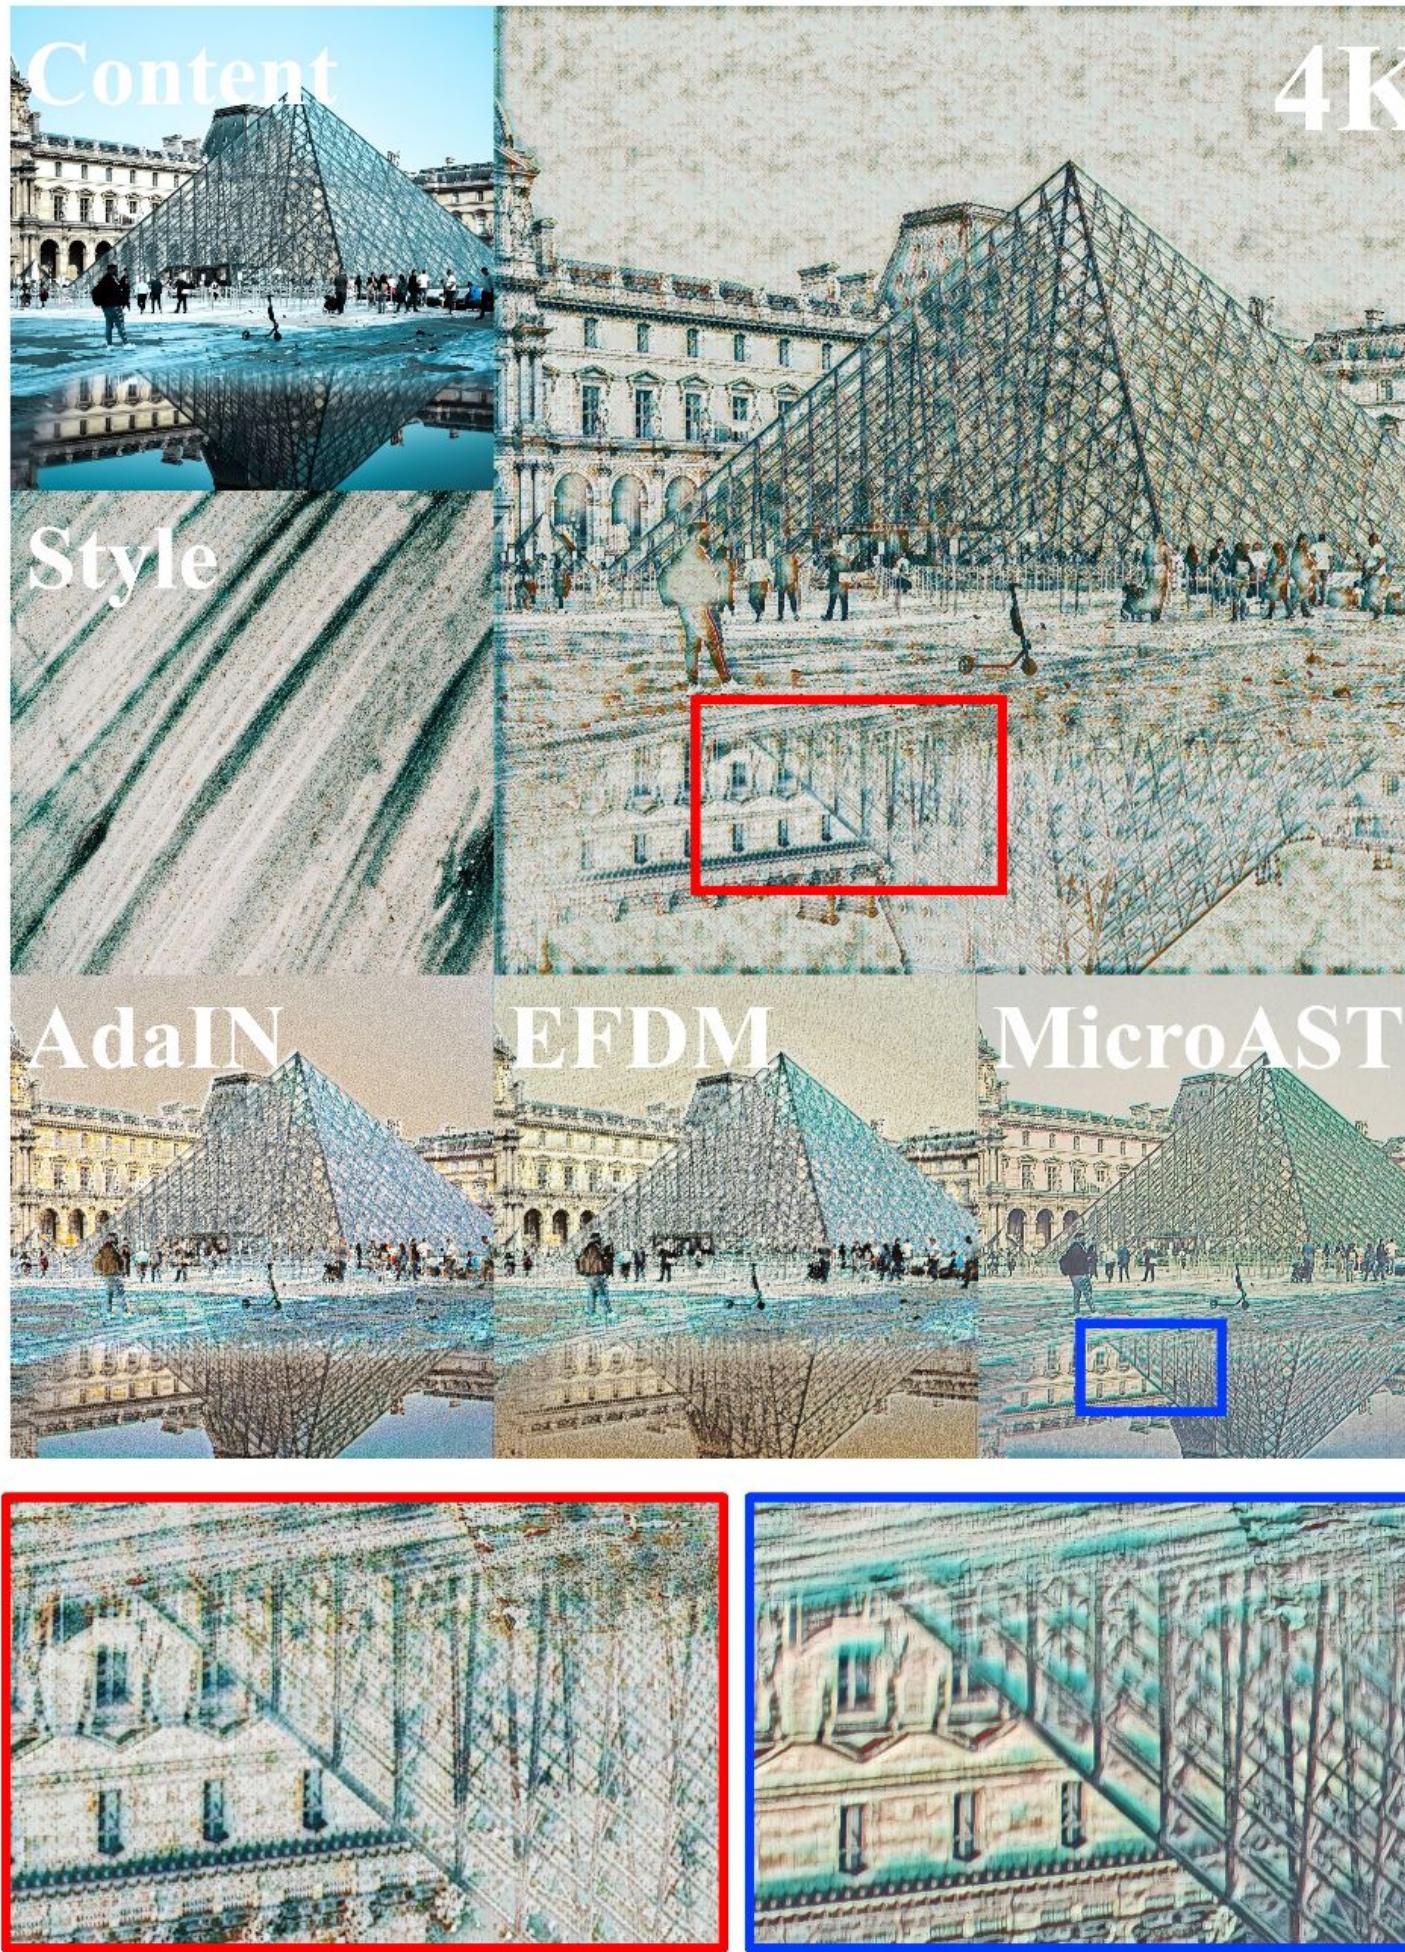
components

Focuses on **hard negatives** for efficient training

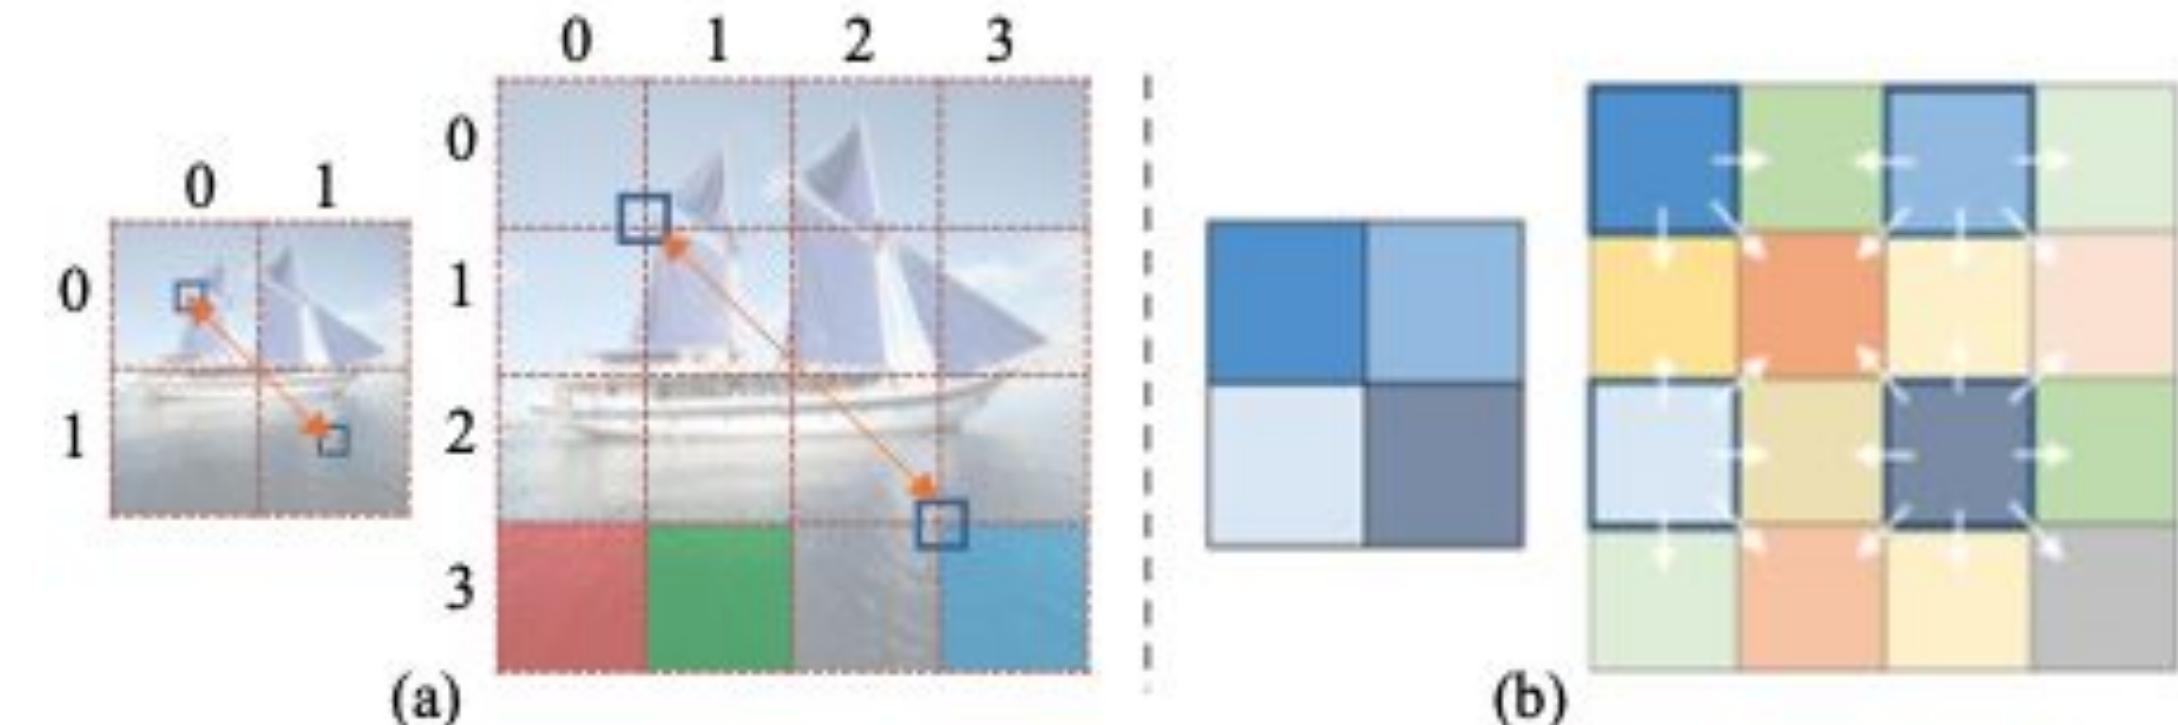
**No pre-trained model at inference**



# 4. NST On Edge (AesFA)



# 5. NST w/ Transformers



## Problem with CNNs:

Limited receptive field → Loses global information

Biased content representation

Loss of fine details in deep networks

## Why Transformers?

Self-attention mechanism captures long-range dependencies

Holistic understanding at each layer

Strong content preservation & style adaptation

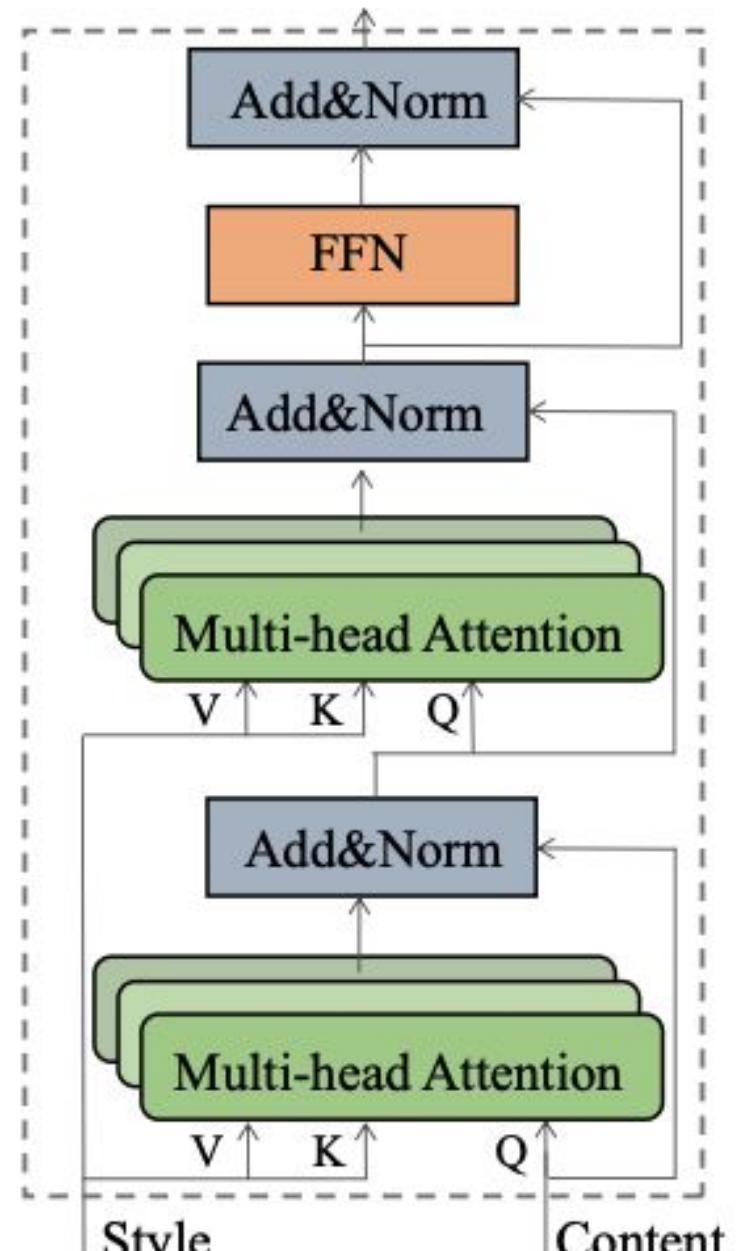
## Key Contributions:

Dual Transformer Encoders: Separate style & content feature extraction

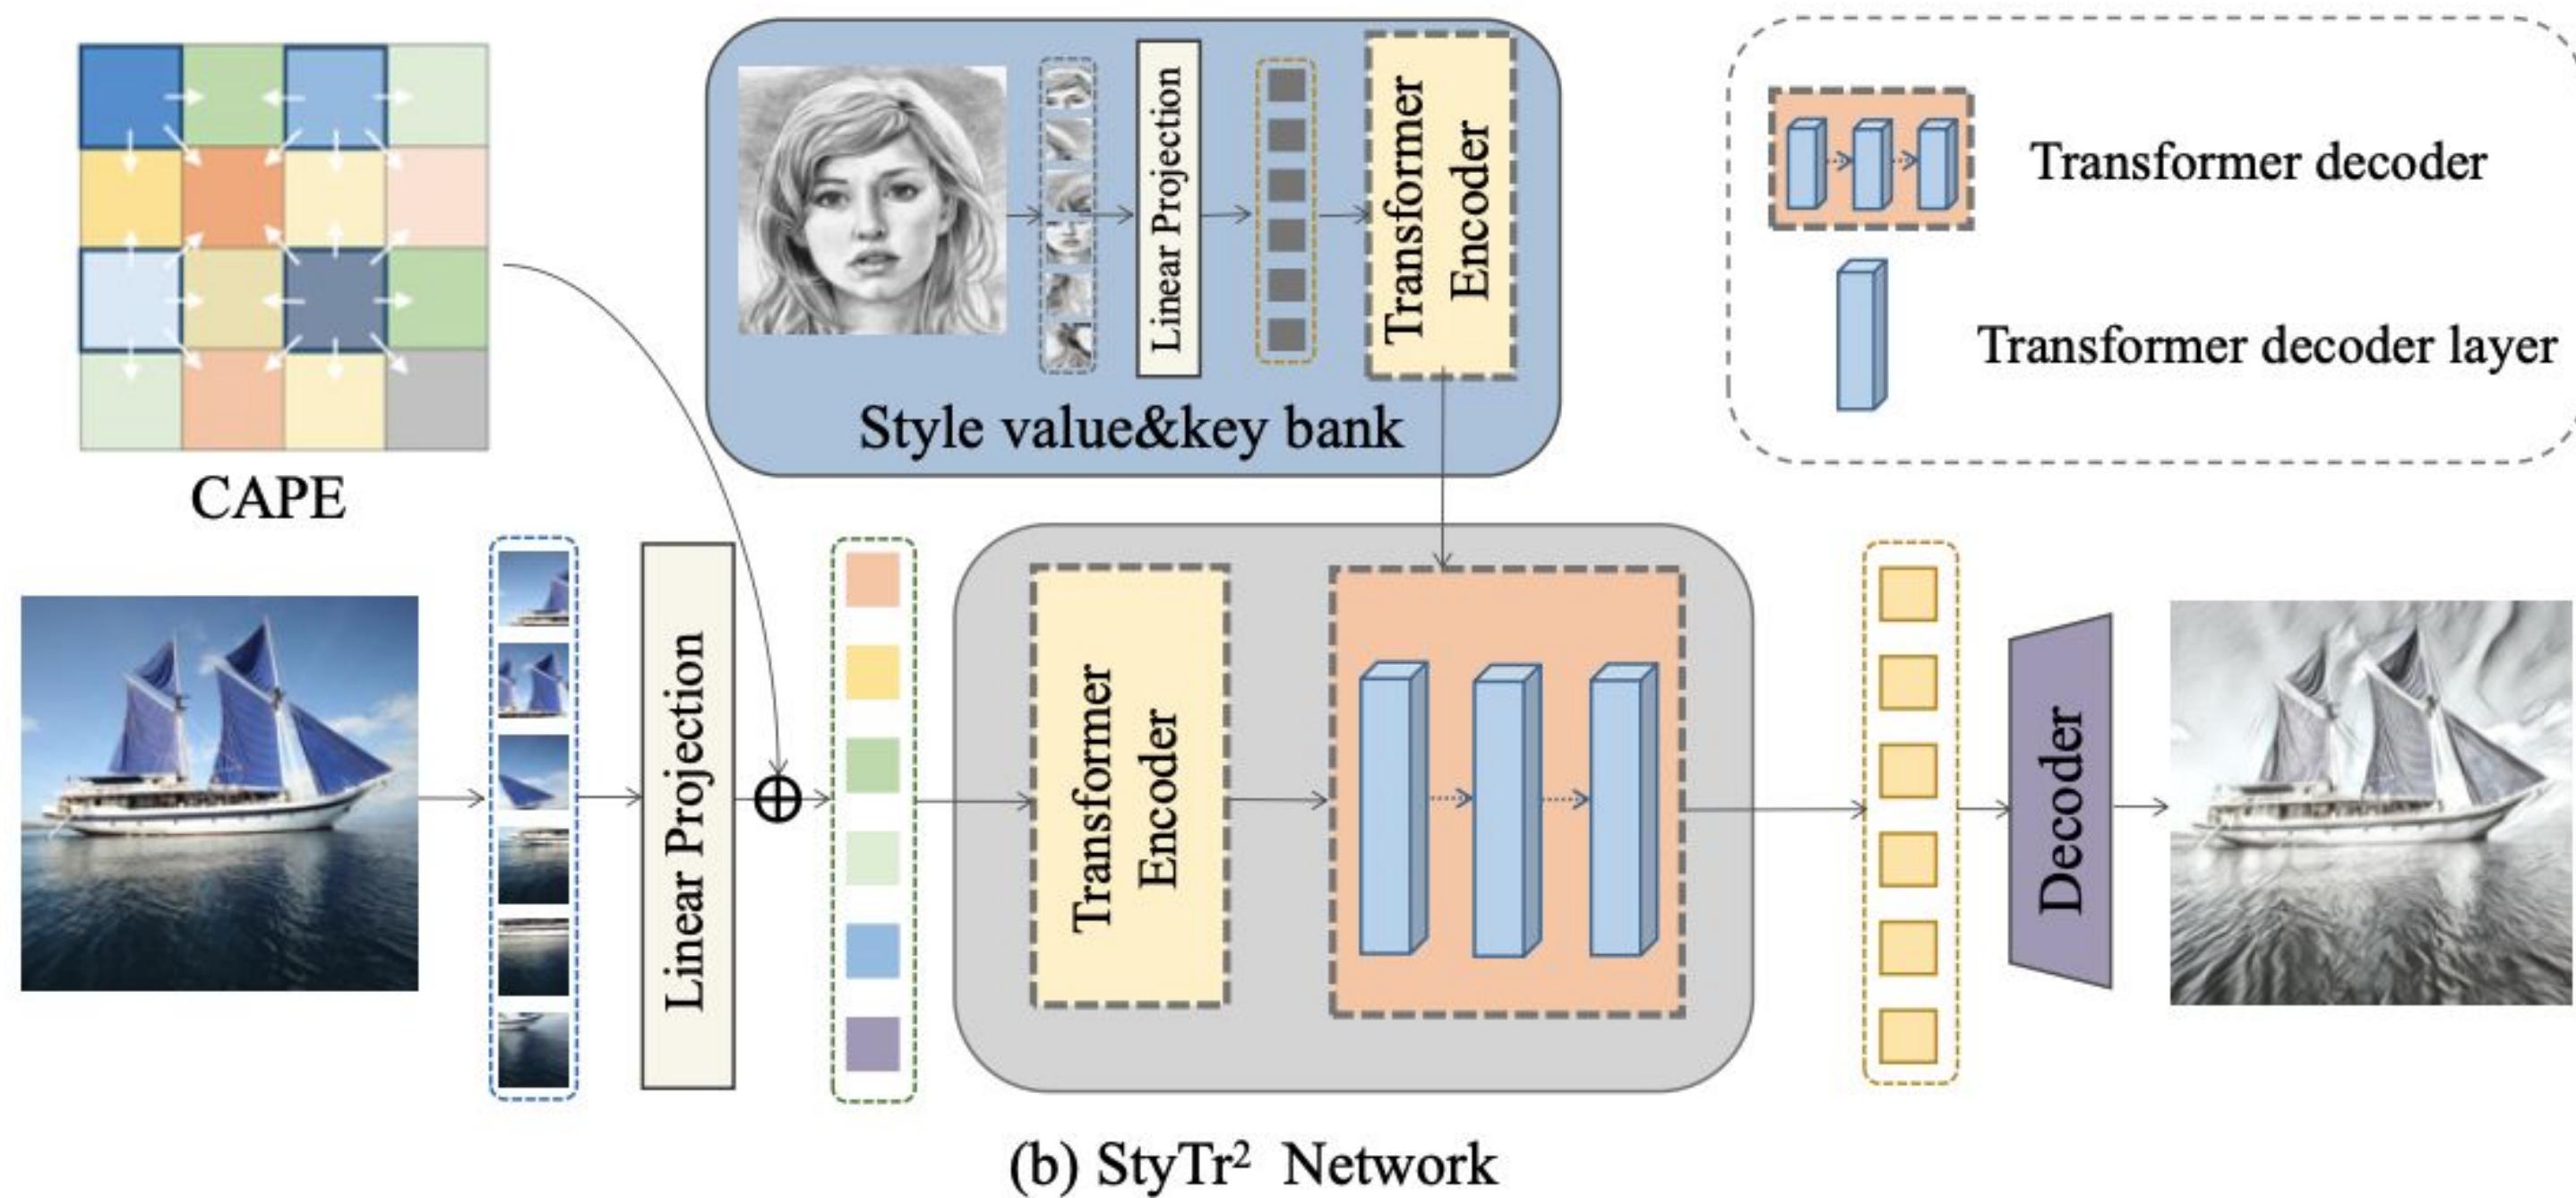
Transformer Decoder: Stylizes content sequence progressively

Content-Aware Positional Encoding (CAPE): Scale-invariant and adjusts to image sizes

# 5. NST w/ Transformers



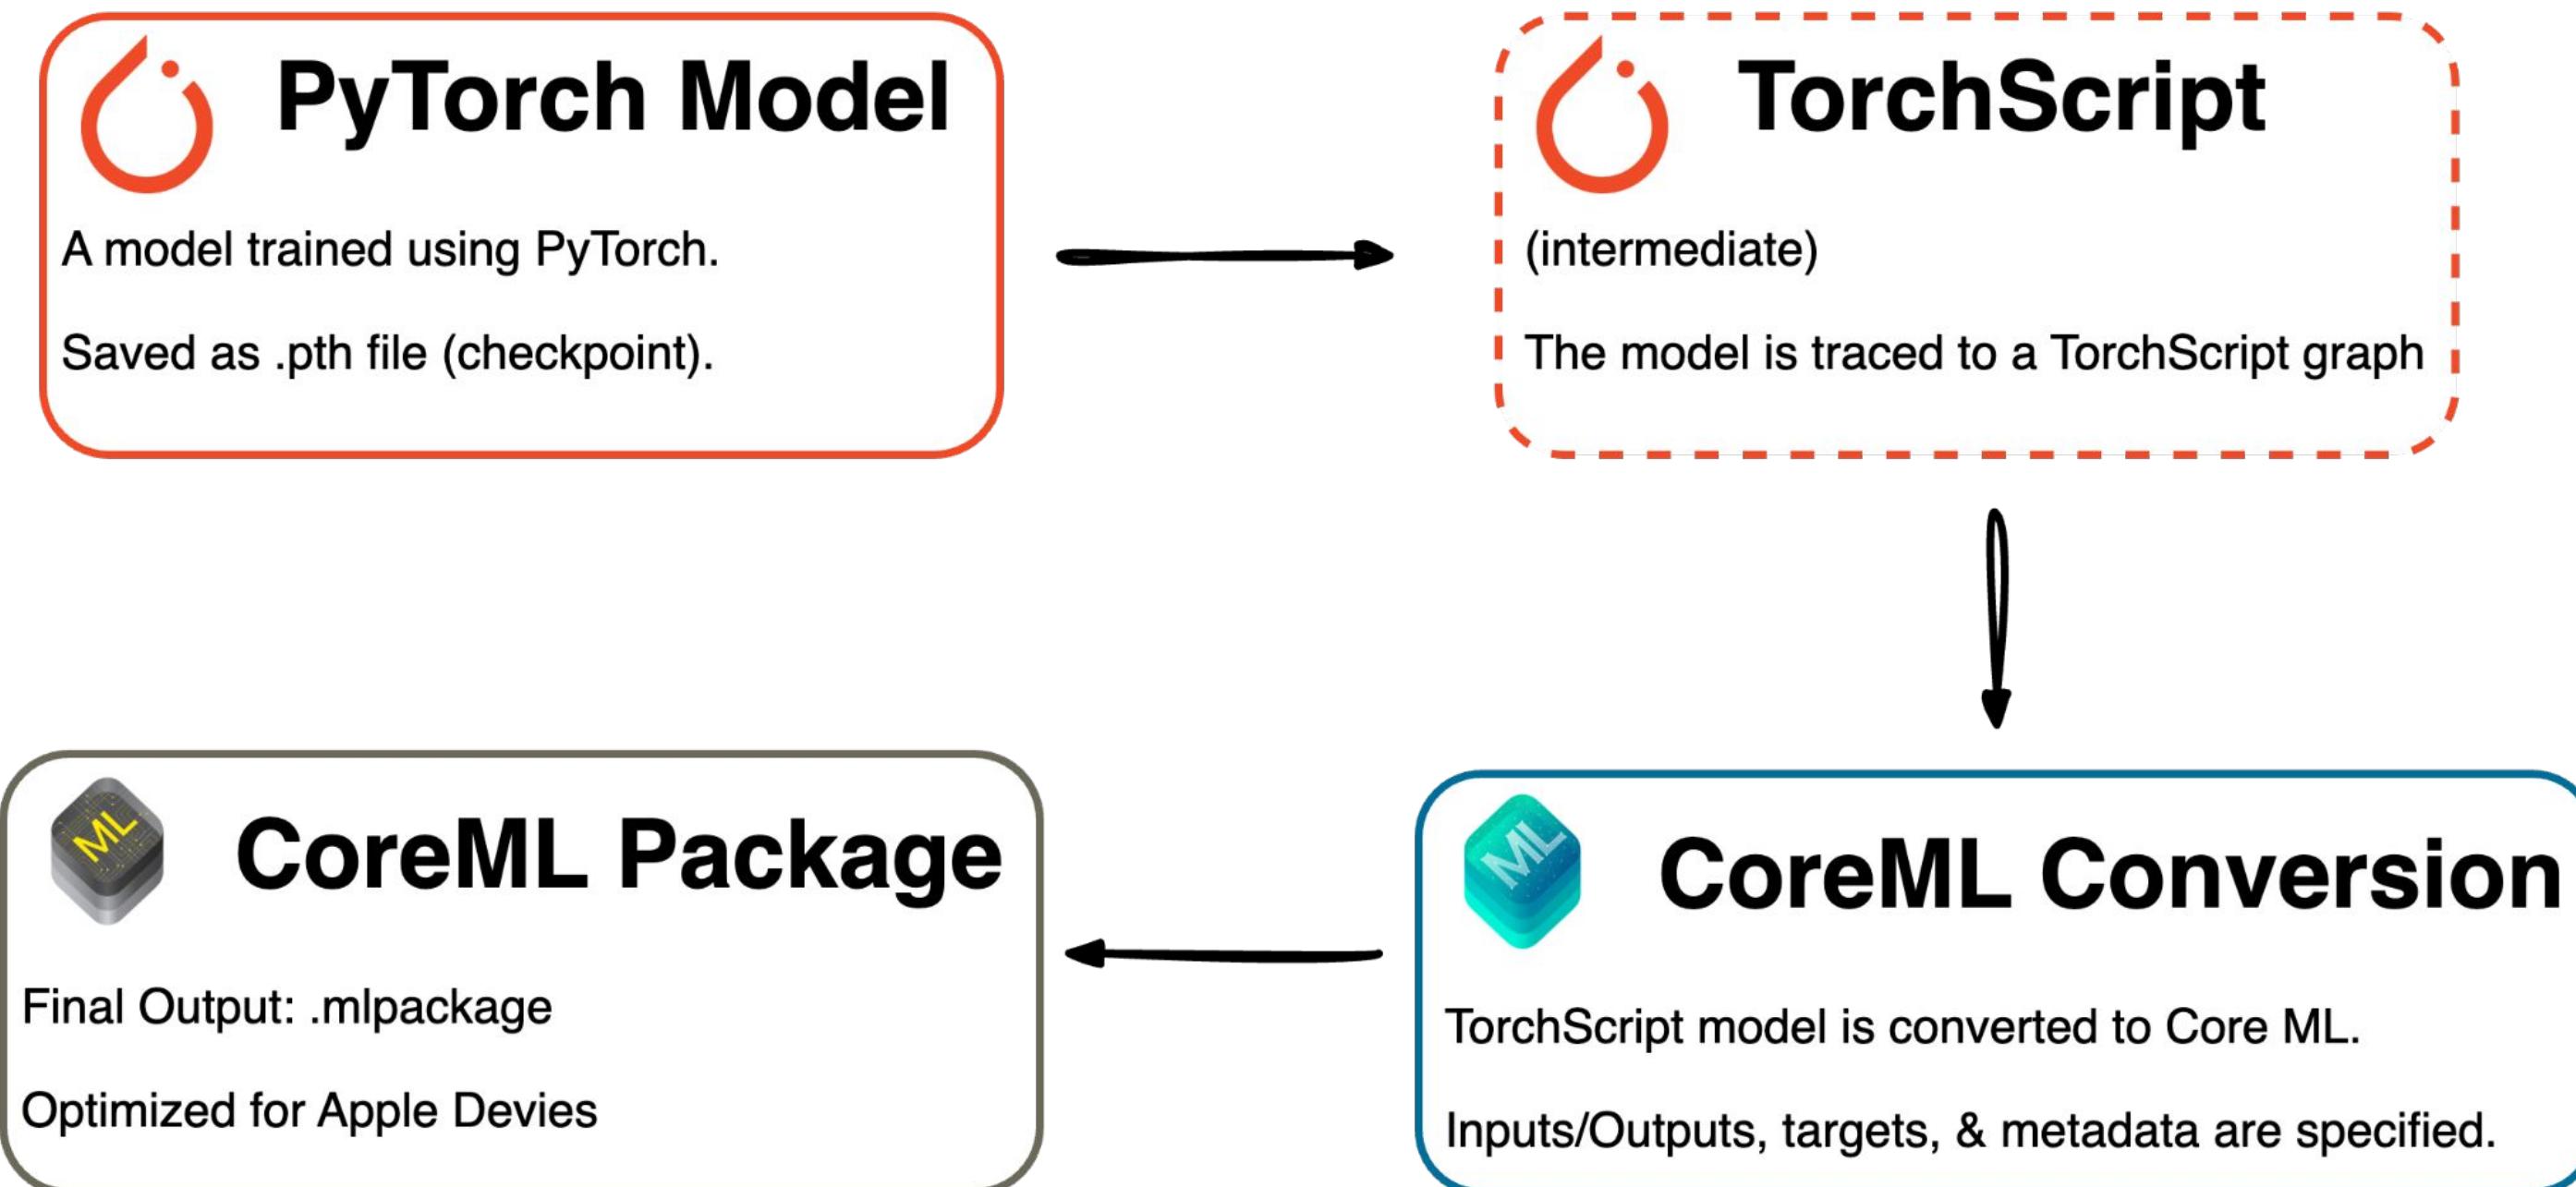
(a) Transformer decoder layer



(b) StyTr<sup>2</sup> Network

# Methodology

# Workflow



# Model-Specific Adaptations

## MobileNet\_AdalN

Innovation: (Encoder → AdalN → Decoder)

Uses **MobileNetV2** (up to layer 13), because it's lightweight and mobile-friendly, instead of VGG-based encoders for feature extraction

Implements AdalN

Uses a separate VGG19-based **perceptual loss** network for:

Content loss: comparing relu4\_1 features

Style loss: Gram matrix comparison of relu1\_1 to relu3\_1 layers

Training: ~4 hours, 20 epochs, 20k images,

## Fast Style Transfer

It was trained in a model-per-style approach and we used pre-trained models for inference

# Model-Specific Adaptations

## StyTr2

**Deprecated implementations** and version-specific-bugs were replaced e.g., `torchvision.ops.misc` package → replaced with hand-coded functions

## AesFA

**Mobile Optimization:** Reworked key parts of the network architecture to make it compatible and efficient for running directly on mobile devices.

**Model & Loss Improvements:** Made internal adjustments to the model's layers and how it calculates style similarity (loss) for better overall performance and stability.

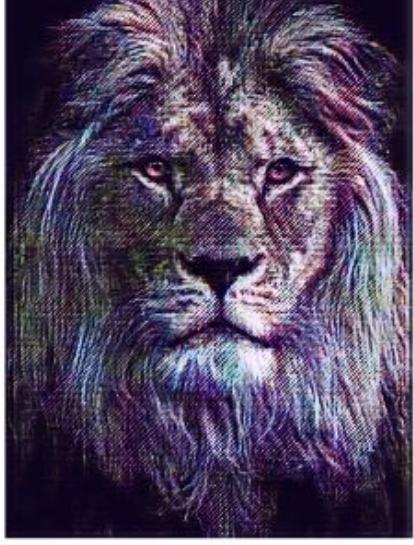
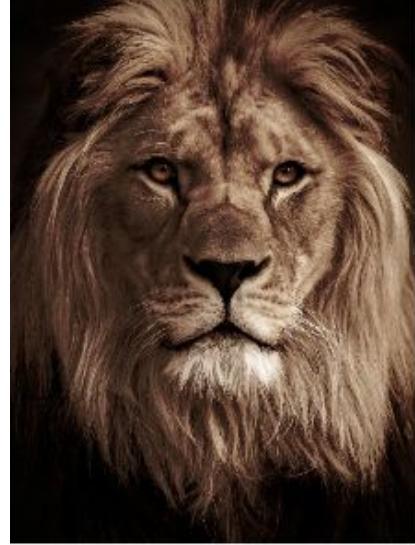
# Results

## Contents

## AdaIN

Laptop

Mobile

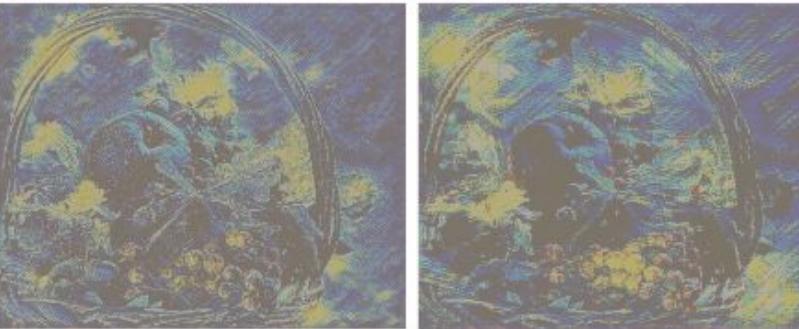


← Style

## FST

Laptop

Mobile



## AesFA

Laptop

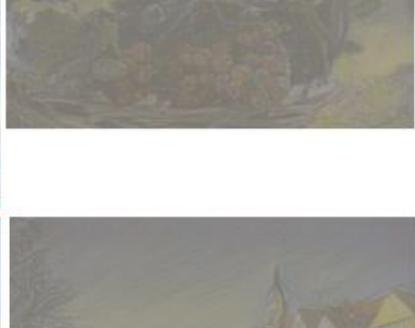
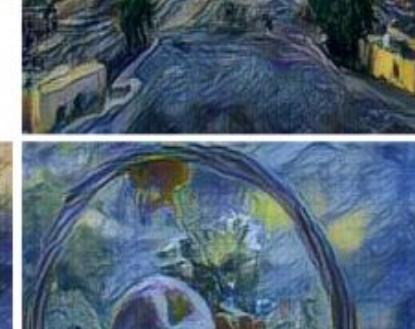
Mobile



## StyTr2

Laptop

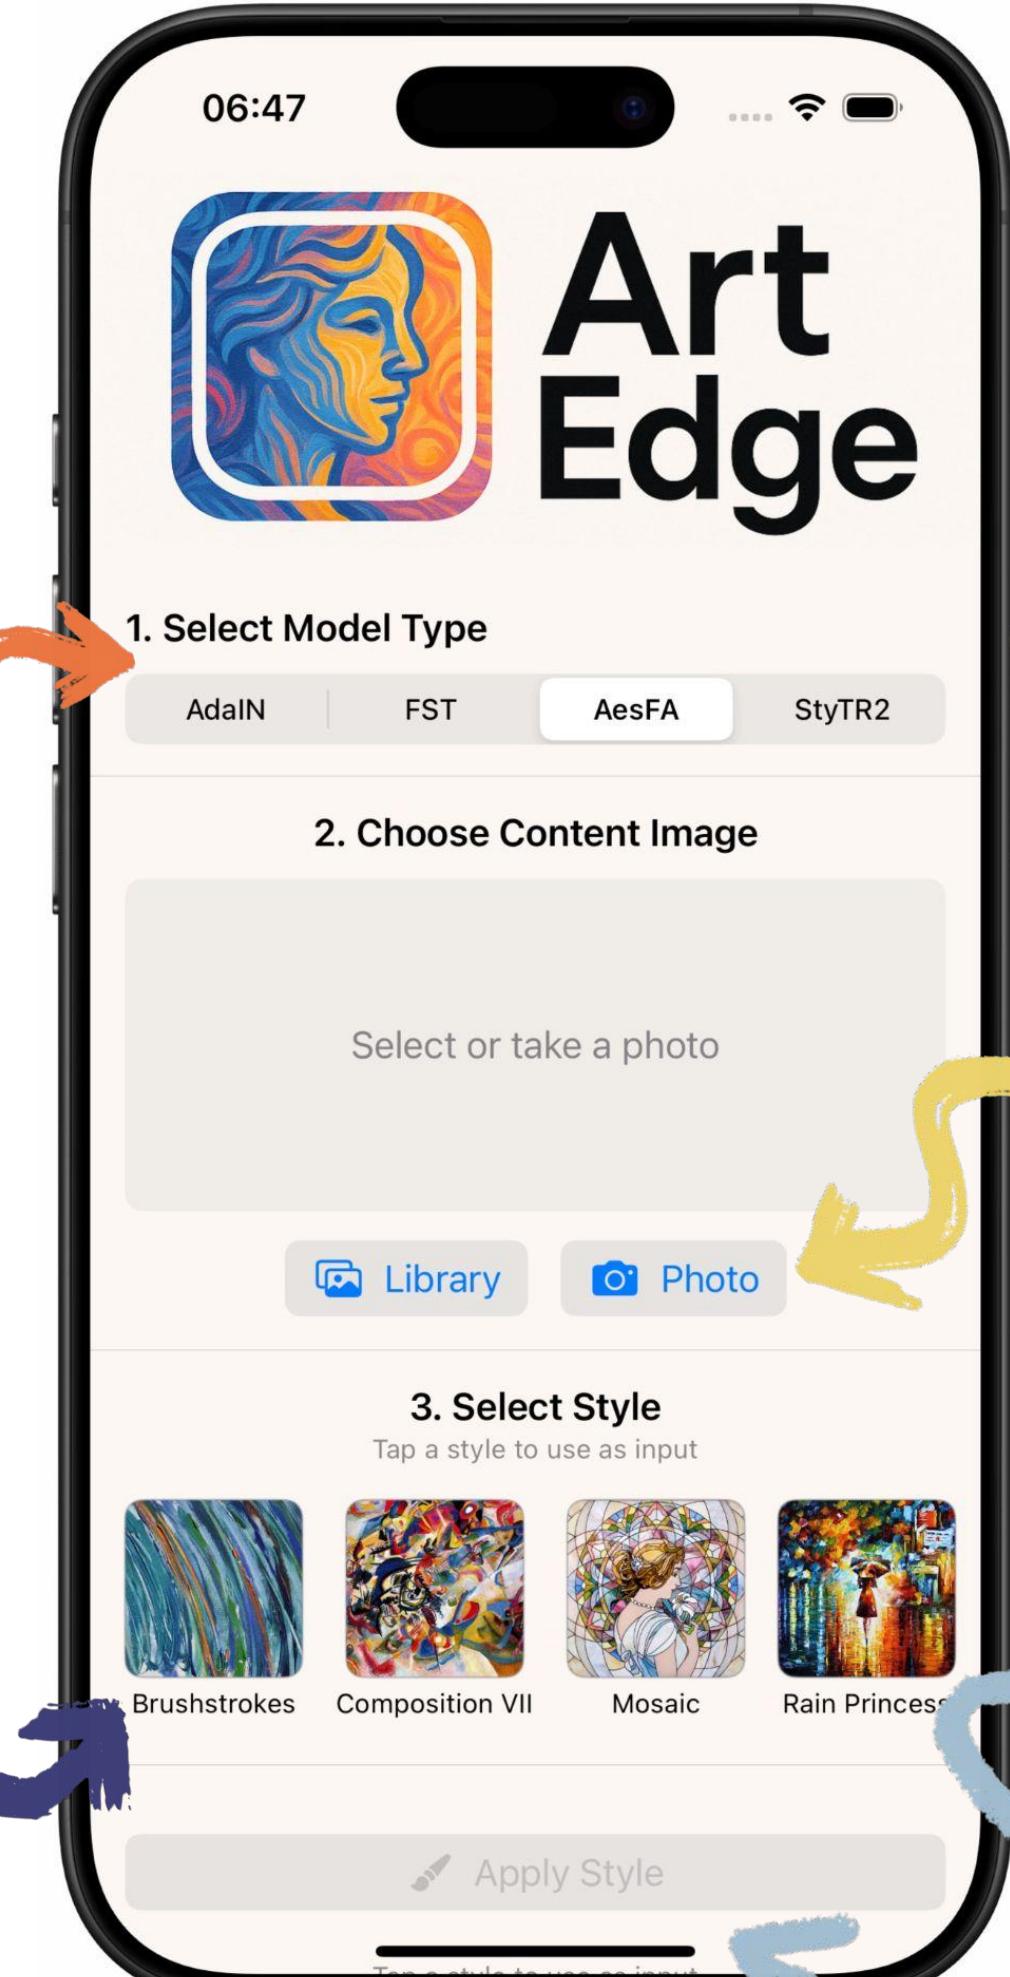
Mobile



# DEMO

Select Different Models on Runtime

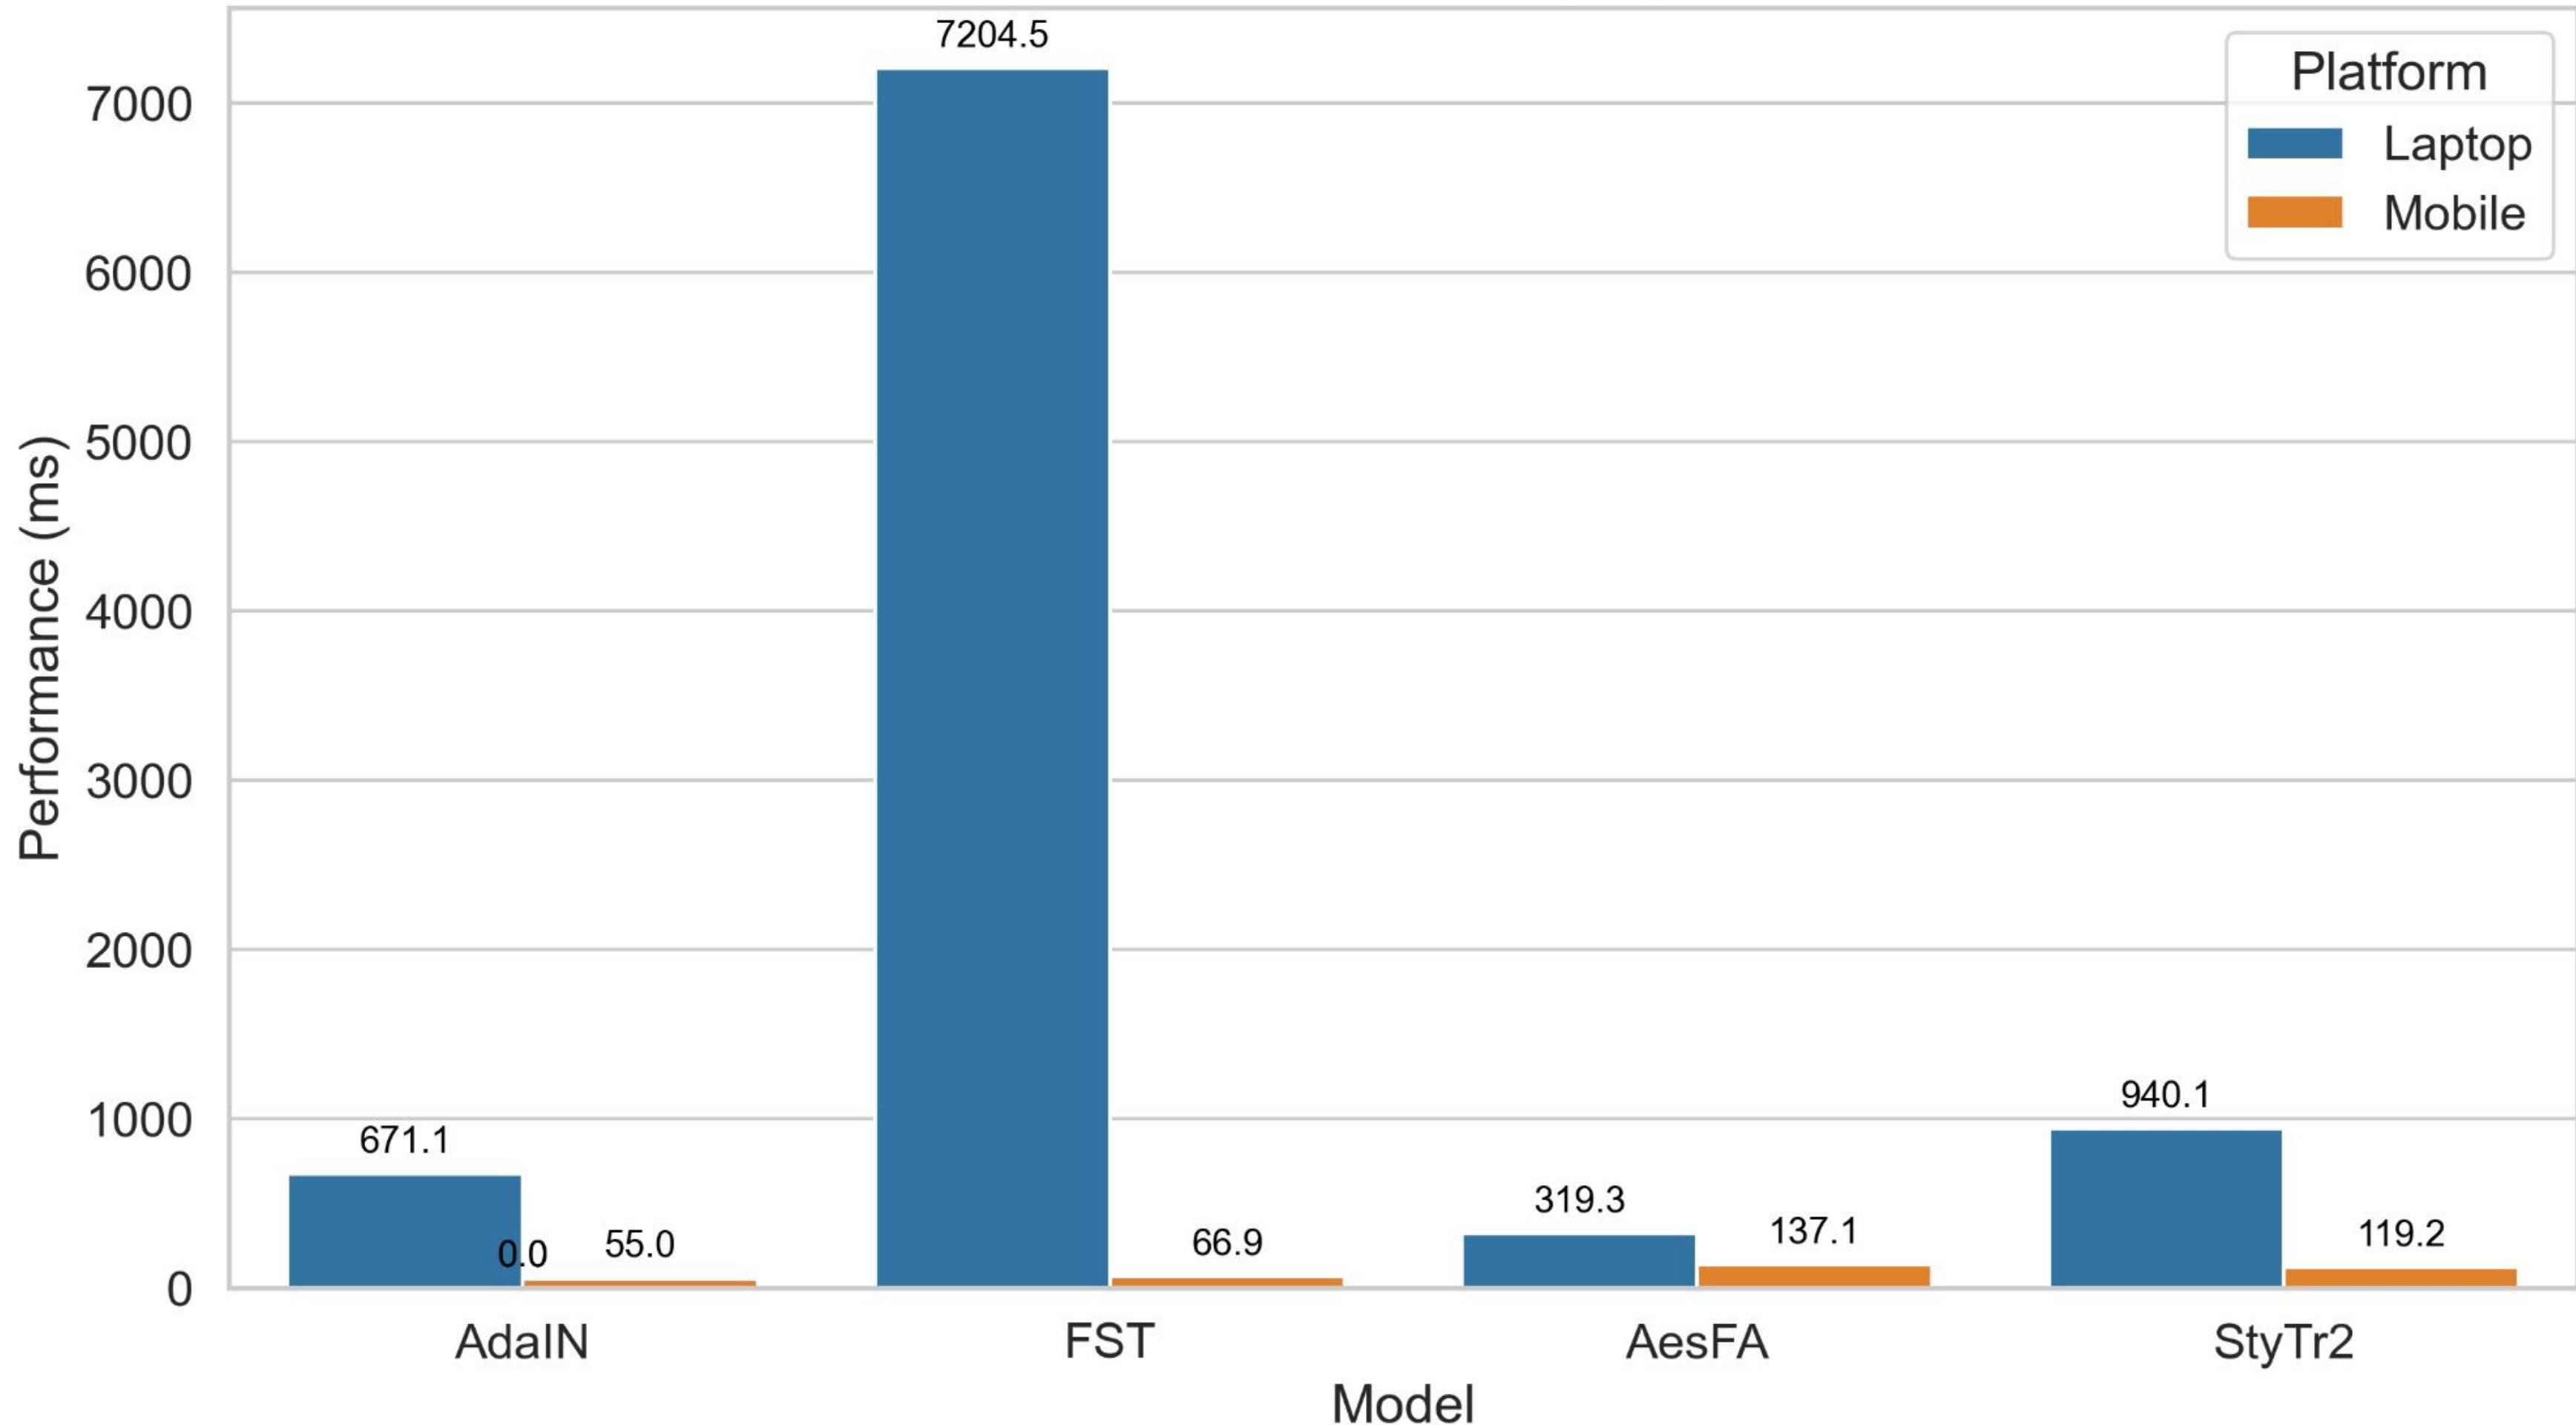
Choose From Pre-Selected Styles



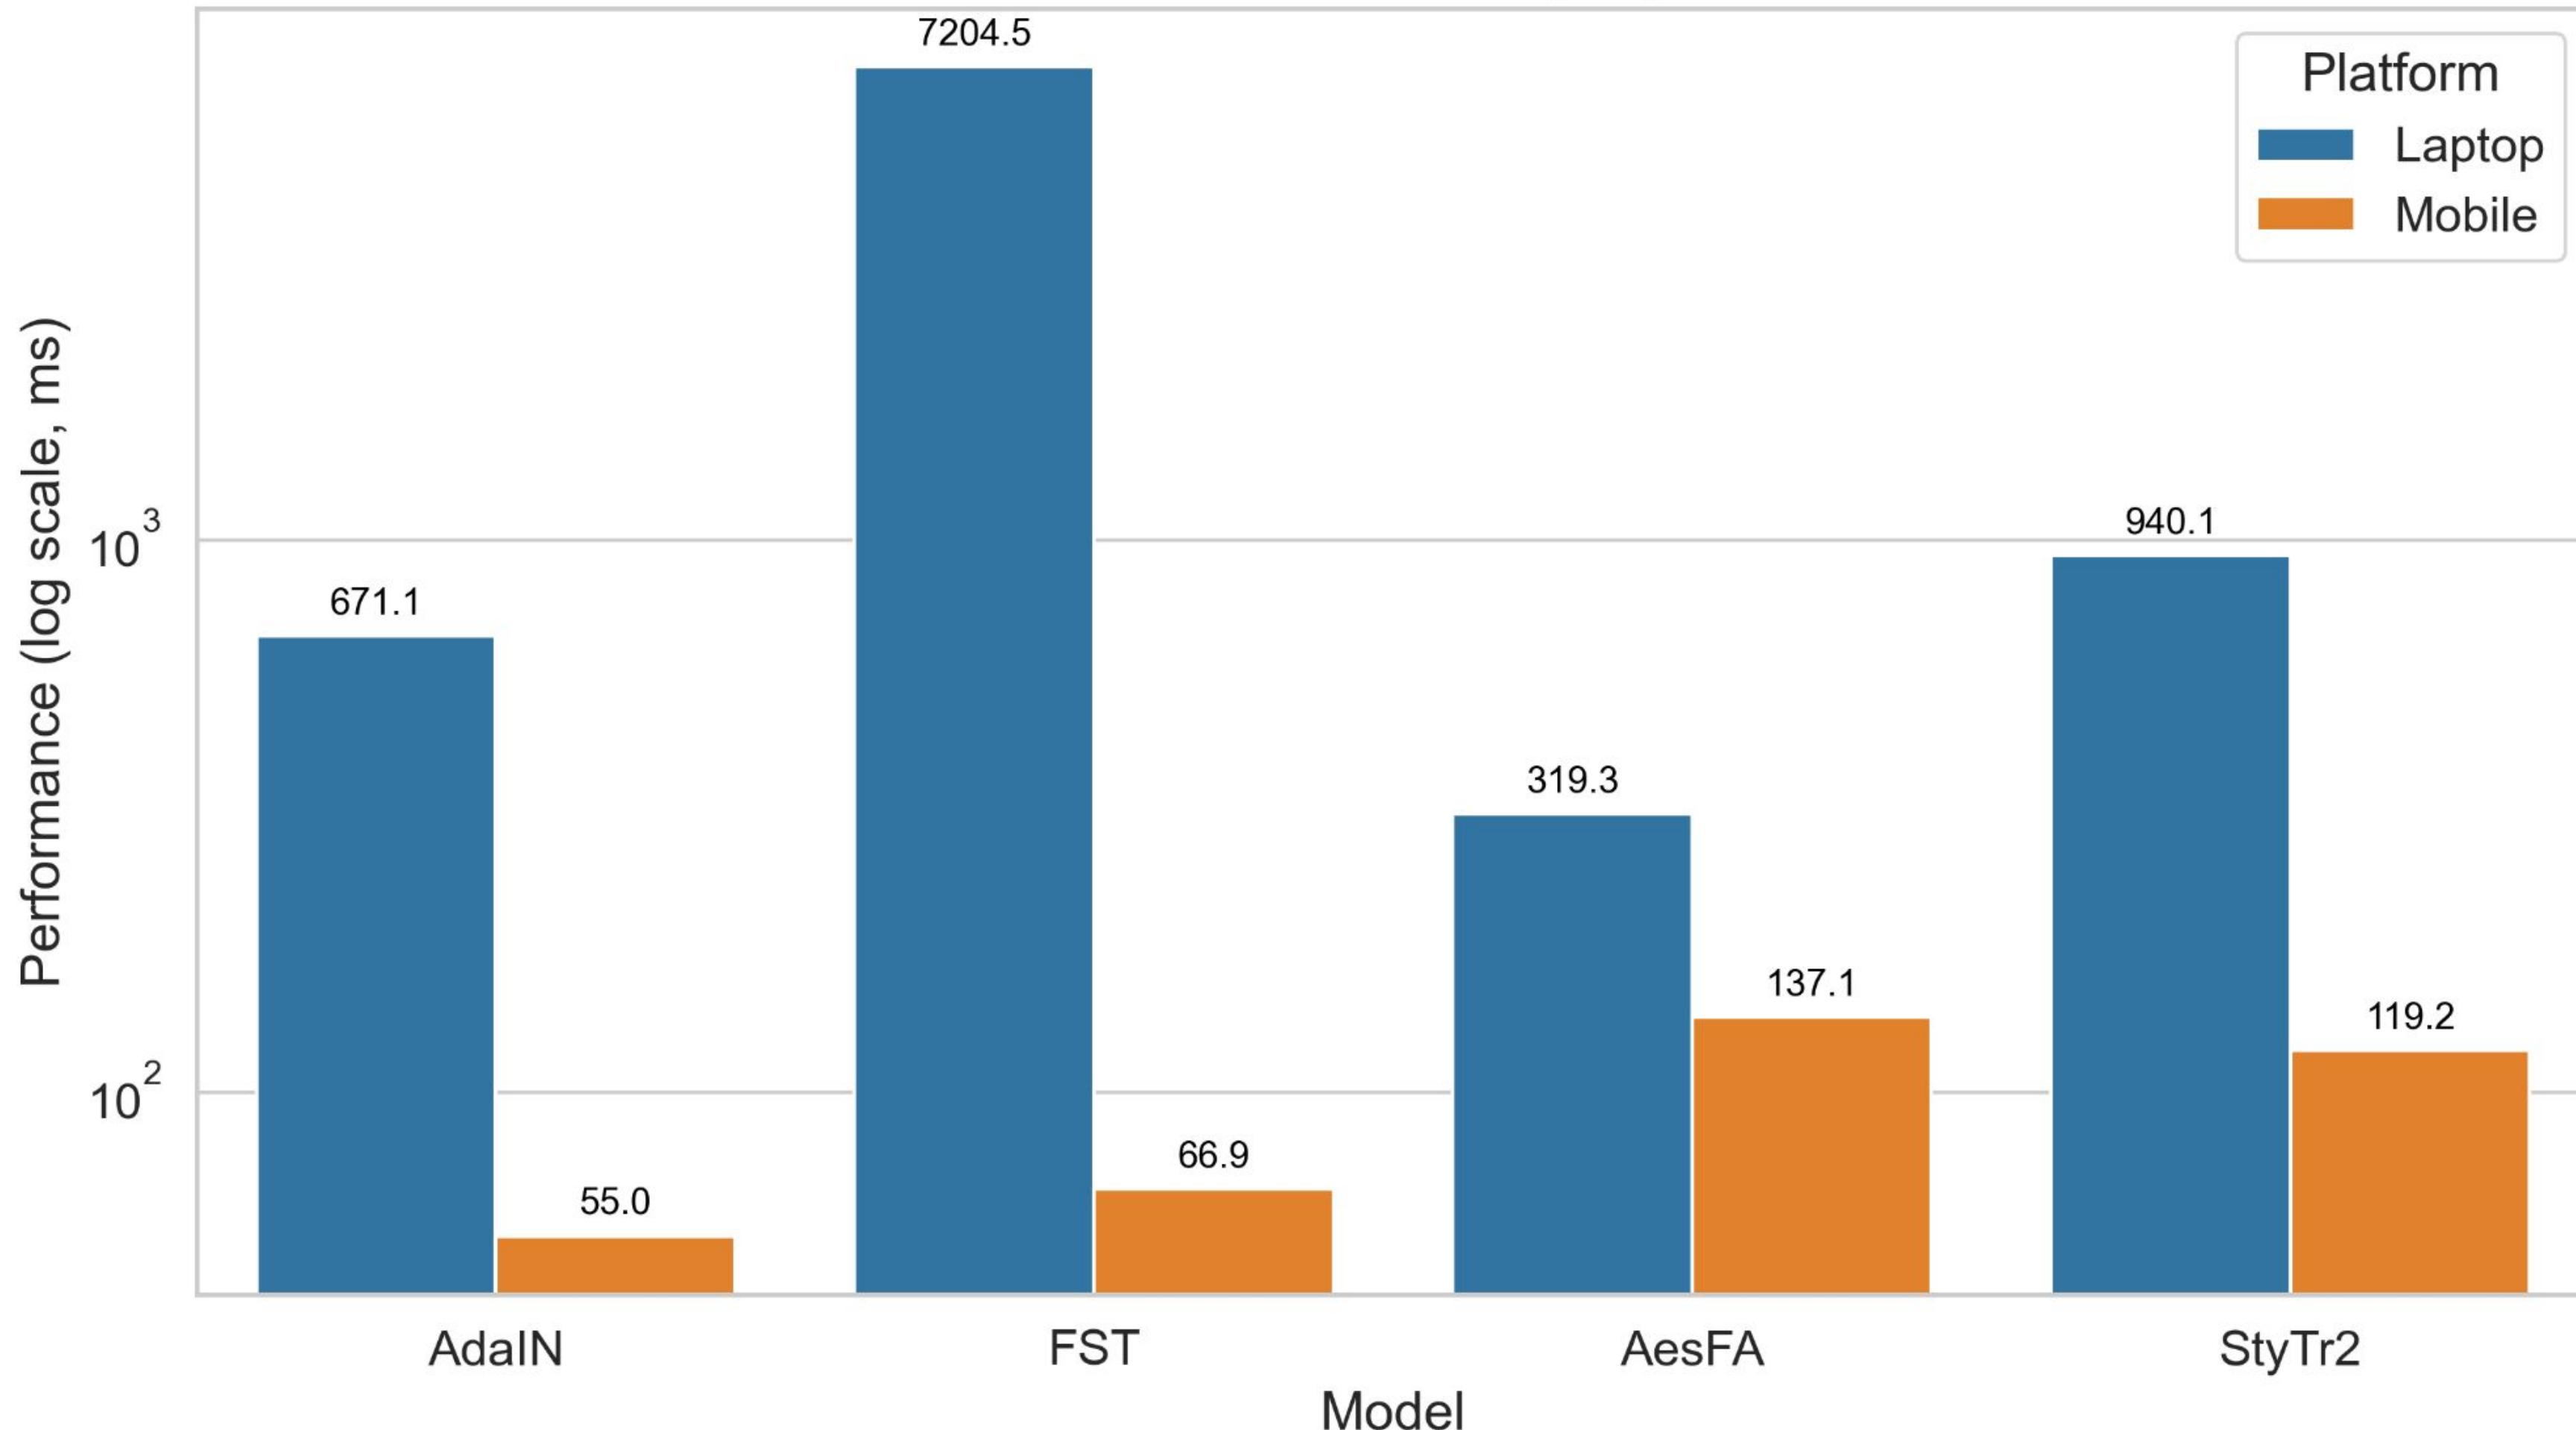
Choose From Your Phone Library or Take a Photo

Apply Style Transfer!

# Model Performance on Laptop vs Mobile



# Model Performance on Laptop vs Mobile



# Conclusion

Integrated four NST models: **MobileNet AdalN**, **Fast Style Transfer**, **AesFA**, and **StyleTr<sup>2</sup>** into a mobile deployment pipeline.

On laptop, all models produce high-quality stylized images, visually faithful to the target style.

On mobile, **Fast Style Transfer** and **AesFA** deliver decent results with good style retention and for speed AesFA is better than FST.

**AesFA** offers the most visually appealing output. **StyleTr<sup>2</sup>** is close but suffers from **lighting artifacts on mobile**, making results appear faded.

Future work will focus on **lighting normalization** and mobile-specific fine-tuning to enhance cross-device visual consistency.

# Thank You!

SlidesCarnival  
for the presentation template

Pexels  
for the art photos